


Linearizability via Order Extension Theorems

Ana Sokolova  UNIVERSITY
of SALZBURG

Dagstuhl, 25.5.2018

- Part I: Concurrent data structures
correctness and performance



- Part II: Order extension results for
verifying linearizability

Concurrent Data Structures

Correctness and Relaxations




Hannes Payer
Google



Tom Henzinger
IST AUSTRIA



Christoph Kirsch
UNIVERSITY
of SALZBURG



Ali Sezgin
UNIVERSITY OF
CAMBRIDGE



Andreas Haas
Google



Michael Lippautz




Andreas Holzer
Google




Helmut Veith
TU
WIEN

Data structures


- Queue FIFO



- Stack LIFO




- Pool unordered




Concurrent data structures


- Queue FIFO



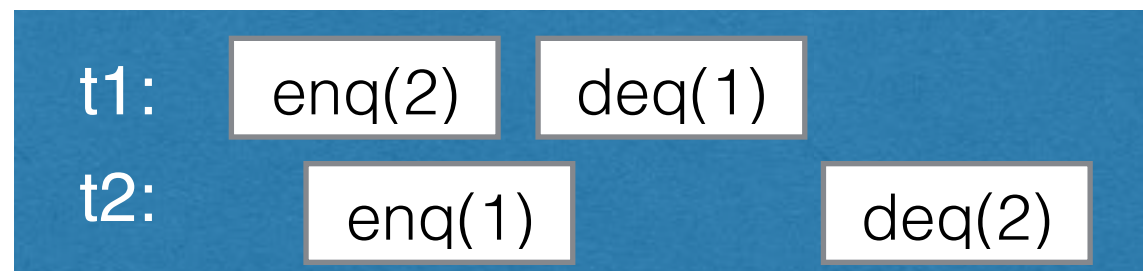
- Stack LIFO



- Pool unordered



Semantics of concurrent data structures



e.g. queues

- Sequential specification = set of legal sequences

e.g. queue legal sequence
enq(1)enq(2)deq(1)deq(2)

- Consistency condition = e.g. linearizability / sequential consistency

e.g. the concurrent history above is a linearizable queue concurrent history

Consistency conditions

there exists a legal sequence that preserves precedence order

Linearizability [Herlihy, Wing '90]

consistency is about extending partial orders to total orders




t1: enq(2)² — deq(1)³
t2: ¹enq(1) — deq(2)⁴

Sequential Consistency [Lamport'79]

there exists a legal sequence that preserves per-thread precedence (program order)

t1: ¹enq(1) — deq(2)⁴
t2: deq(1)² — enq(2)³

Performance and scalability



Relaxations allow trading

correctness
for
performance

provide the **potential**
for better-performing
implementations

Relaxing the Semantics


Quantitative relaxations
Henzinger, Kirsch, Payer, Sezgin, S. POPL13

- Sequential specification = set of legal sequences
- Consistency condition = e.g. linearizability / sequential consistency


Local linearizability
Haas, Henzinger, Holzer, ..., S, Veith CONCUR16

Lead to scalable implementations

e.g. k-FIFO, k-Stack




locally linearizable distributed implementation



local inserts / global removes


LLD Φ
LL+D Φ

Performance



(a) Queues, LL queues, and “queue-like” pools


Performance



LLD Φ performs significantly better than Φ

(a) Queues, LL queues, and “queue-like” pools

Performance



(a) Queues, LL queues, and “queue-like” pools

Linearizability via Order Extension Theorems

joint work with



Harald Woracek



foundational results
for
verifying linearizability

Inspiration

As well as
Reducing Linearizability to
State Reachability
[Bouajjani, Emmi, Enea, Hamza]
ICALP15 + ...

Queue sequential specification (axiomatic)

s is a legal queue sequence
iff

1. **s** is a legal pool sequence, and
2. $\text{enq}(x) <_{\mathbf{s}} \text{enq}(y) \wedge \text{deq}(y) \in \mathbf{s} \Rightarrow \text{deq}(x) \in \mathbf{s} \wedge \text{deq}(x) <_{\mathbf{s}} \text{deq}(y)$

Queue linearizability (axiomatic)

Henzinger, Sezgin, Vafeiadis CONCUR13

h is queue linearizable
iff

1. **h** is pool linearizable, and
2. $\text{enq}(x) <_{\mathbf{h}} \text{enq}(y) \wedge \text{deq}(y) \in \mathbf{h} \Rightarrow \text{deq}(x) \in \mathbf{h} \wedge \text{deq}(y) \not<_{\mathbf{h}} \text{deq}(x)$

precedence order

Concurrent Queues

Data independence \Rightarrow verifying executions where each value is enqueued at most once is sound

Reduction to **assertion checking** = exclusion of "bad patterns"

Value v dequeued without being enqueued

$\text{deq} \Rightarrow v$



Value v dequeued before being enqueued

$\text{deq} \Rightarrow v$ $\text{enq}(v)$




Value v dequeued twice

$\text{deq} \Rightarrow v$ $\text{deq} \Rightarrow v$



Value v_1 and v_2 dequeued in the wrong order

$\text{enq}(v_1)$ $\text{enq}(v_2)$ $\text{deq} \Rightarrow v_2$ $\text{deq} \Rightarrow v_1$



Concurrent Queues

Data independence \Rightarrow verifying executions where each value is enqueued at most once is sound

Reduction to **assertion checking** = exclusion of "bad patterns"

Value v dequeued without being enqueued

$\text{deq} \Rightarrow v$




Value v dequeued before being enqueued

$\text{deq} \Rightarrow v$ $\text{enq}(v)$




Value v dequeued twice

$\text{deq} \Rightarrow v$ $\text{deq} \Rightarrow v$




Value v_1 and v_2 dequeued in the wrong order

$\text{enq}(v_1)$ $\text{enq}(v_2)$ $\text{deq} \Rightarrow v_2$ $\text{deq} \Rightarrow v_1$



Dequeue wrongly returns empty

$\text{deq} \Rightarrow \text{empty}$



Linearizability verification

Data structure

- signature Σ - set of method calls including data values
- sequential specification $S \subseteq \Sigma^*$, prefix closed

identify sequences with total orders

Sequential specification via violations

Extract a set of violations V , relations on Σ , such that $\mathbf{s} \in S$ iff \mathbf{s} has no violations

it is easy to find a large CV,
but difficult to find a small representative

$$\mathcal{P}(\mathbf{s}) \cap V = \emptyset$$

Linearizability verification

Find a set of violations CV such that: every interval order with no CV violations extends to a total order with no V violations.

we build
CV iteratively
from V

legal sequence

concurrent history

Pool without empty removals

Pool sequential specification (axiomatic)

s is a legal pool (without empty removals) sequence

iff
1. $\text{rem}(x) \in \mathbf{s} \Rightarrow \text{ins}(x) \in \mathbf{s} \wedge \text{ins}(x) <_{\mathbf{s}} \text{rem}(x)$

V violations
 $\text{rem}(x) <_{\mathbf{s}} \text{ins}(x)$

Pool linearizability (axiomatic)

h is pool (without empty removals) linearizable

iff
1. $\text{rem}(x) \in \mathbf{h} \Rightarrow \text{ins}(x) \in \mathbf{h} \wedge \text{rem}(x) \not\prec_{\mathbf{h}} \text{ins}(x)$

CV violations
= V violations

Queue without empty removals

Queue sequential specification (axiomatic)

s is a legal queue (without empty removals) sequence
iff

1. $\text{deq}(x) \in \mathbf{s} \Rightarrow \text{enq}(x) \in \mathbf{s} \wedge \text{enq}(x) <_{\mathbf{s}} \text{deq}(x)$
2. $\text{enq}(x) <_{\mathbf{s}} \text{enq}(y) \wedge \text{deq}(y) \in \mathbf{s} \Rightarrow \text{deq}(x) \in \mathbf{s} \wedge \text{deq}(x) <_{\mathbf{s}} \text{deq}(y)$

V violations
 $\text{deq}(x) <_{\mathbf{s}} \text{enq}(x)$
and
 $\text{enq}(x) <_{\mathbf{s}} \text{enq}(y) \wedge$
 $\text{deq}(y) <_{\mathbf{s}} \text{deq}(x)$

Queue linearizability (axiomatic)

h is queue (without empty removals) linearizable
iff

1. $\text{rem}(x) \in \mathbf{h} \Rightarrow \text{ins}(x) \in \mathbf{h} \wedge \text{rem}(x) \not<_{\mathbf{h}} \text{ins}(x)$
2. $\text{enq}(x) <_{\mathbf{h}} \text{enq}(y) \wedge \text{deq}(y) \in \mathbf{h} \Rightarrow \text{deq}(x) \in \mathbf{h} \wedge \text{deq}(y) \not<_{\mathbf{h}} \text{deq}(x)$

CV violations
= V violations

Pool

infinite
inductive
violations

Pool sequential specification (axiomatic)

s is a legal pool (with empty removals) sequence
iff

1. $\text{rem}(x) \in \mathbf{s} \Rightarrow \text{ins}(x) \in \mathbf{s} \wedge \text{ins}(x) <_{\mathbf{s}} \text{rem}(x)$
2. $\text{rem}(\perp) <_{\mathbf{s}} \text{rem}(x) \Rightarrow \text{rem}(\perp) <_{\mathbf{s}} \text{ins}(x) \wedge \text{ins}(x) <_{\mathbf{s}} \text{rem}(\perp) \Rightarrow \text{rem}(x) <_{\mathbf{s}} \text{rem}(\perp)$

\forall violations
 $\text{rem}(x) <_{\mathbf{s}} \text{ins}(x)$
and
 $\text{ins}(x) <_{\mathbf{s}} \text{rem}(\perp) <_{\mathbf{s}} \text{rem}(x)$

Pool linearizability (axiomatic)

h is pool (with empty removals) linearizable
iff

1. $\text{rem}(x) \in \mathbf{h} \Rightarrow \text{ins}(x) \in \mathbf{h} \wedge \text{rem}(x) \not<_{\mathbf{h}} \text{ins}(x)$
2.

infinitely many CV violations

$\text{ins}(x_1) <_{\mathbf{h}} \text{rem}(\perp) \wedge \text{ins}(x_2) <_{\mathbf{h}} \text{rem}(x_1) \wedge \dots \wedge \text{ins}(x_{n+1}) <_{\mathbf{h}} \text{rem}(x_n) \wedge \text{rem}(\perp) <_{\mathbf{h}} \text{rem}(x_{n+1})$

It works for

- Pool without empty removals
- Queue without empty removals
- Priority queue without empty removals
- Pool
- Queue
- Priority queue

Thank You !

But not yet for Stack:
infinite CV violations
without clear
inductive structure

Exploring the space of
data structures
as well as new ideas
for problematic cases