

Local Linearizability

Ana Sokolova  UNIVERSITY
of SALZBURG

Local Linearizability

Ana Sokolova  UNIVERSITY
of SALZBURG

joint work with:

Andreas Haas  UNIVERSITY
of SALZBURG

Andreas Holzer  UNIVERSITY OF
TORONTO

Michael Lippautz  UNIVERSITY
of SALZBURG

Ali Sezgin  UNIVERSITY OF
CAMBRIDGE

Tom Henzinger  IST AUSTRIA

Christoph Kirsch  UNIVERSITY
of SALZBURG

Hannes Payer 

Helmut Veith  TU
WIEN

Concurrent Data Structures

Correctness and Performance

Semantics of concurrent data structures

Semantics of concurrent data structures

e.g. pools, queues, stacks

Semantics of concurrent data structures

t1: enq(2) deq(1)
t2: enq(1) deq(2)

e.g. pools, queues, stacks

Semantics of concurrent data structures

t1: enq(2) deq(1)
t2: enq(1) deq(2)

e.g. pools, queues, stacks

- Sequential specification = set of legal sequences
- Consistency condition = e.g. linearizability / sequential consistency

Semantics of concurrent data structures

t1: enq(2) deq(1)
t2: enq(1) deq(2)

e.g. pools, queues, stacks

- Sequential specification = set of legal sequences
 - e.g. queue legal sequence
 $\text{enq}(1)\text{enq}(2)\text{deq}(1)\text{deq}(2)$
- Consistency condition = e.g. linearizability / sequential consistency

Semantics of concurrent data structures

t1: enq(2) deq(1)
t2: enq(1) deq(2)

e.g. pools, queues, stacks

- Sequential specification = set of legal sequences
 - e.g. queue legal sequence
 $\text{enq}(1)\text{enq}(2)\text{deq}(1)\text{deq}(2)$
- Consistency condition = e.g. linearizability / sequential consistency

e.g. the concurrent history above is a linearizable queue concurrent history

Consistency conditions

Linearizability [Herlihy,Wing '90]



Sequential Consistency [Lamport'79]

Consistency conditions

there exists a sequential witness that preserves precedence

Linearizability [Herlihy,Wing '90]



Sequential Consistency [Lamport'79]

Consistency conditions

there exists a sequential witness that preserves precedence

Linearizability [Herlihy,Wing '90]



Sequential Consistency [Lamport'79]

Consistency conditions

there exists a sequential witness that preserves precedence

Linearizability [Herlihy,Wing '90]



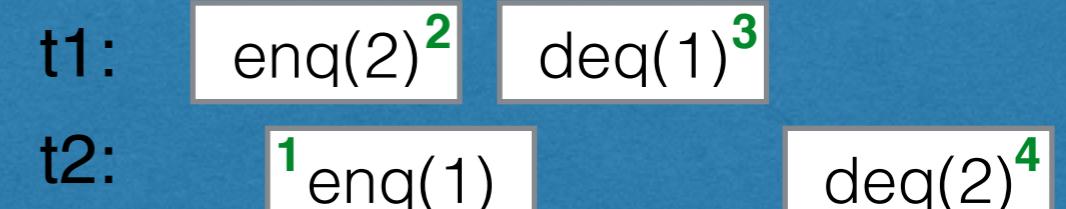
t1:	enq(2) ²	deq(1) ³
t2:	¹ enq(1)	deq(2) ⁴

Sequential Consistency [Lamport'79]

Consistency conditions

there exists a sequential witness that preserves precedence

Linearizability [Herlihy,Wing '90]



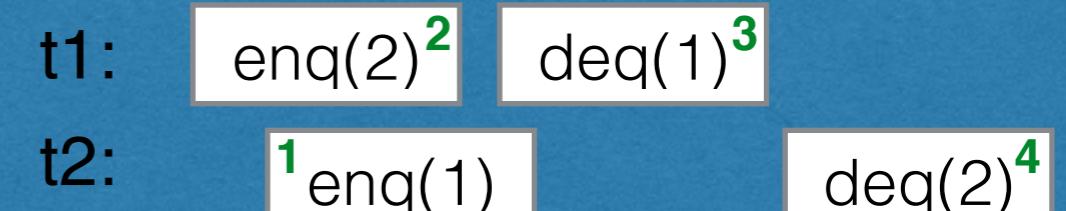
Sequential Consistency [Lamport'79]

there exists a sequential witness that preserves per-thread precedence (program order)

Consistency conditions

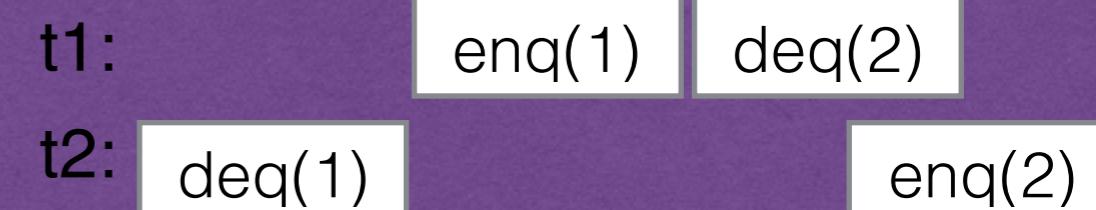
there exists a sequential witness that preserves precedence

Linearizability [Herlihy,Wing '90]



Sequential Consistency [Lamport'79]

there exists a sequential witness that preserves per-thread precedence (program order)



Consistency conditions

there exists a sequential witness that preserves precedence

Linearizability [Herlihy,Wing '90]



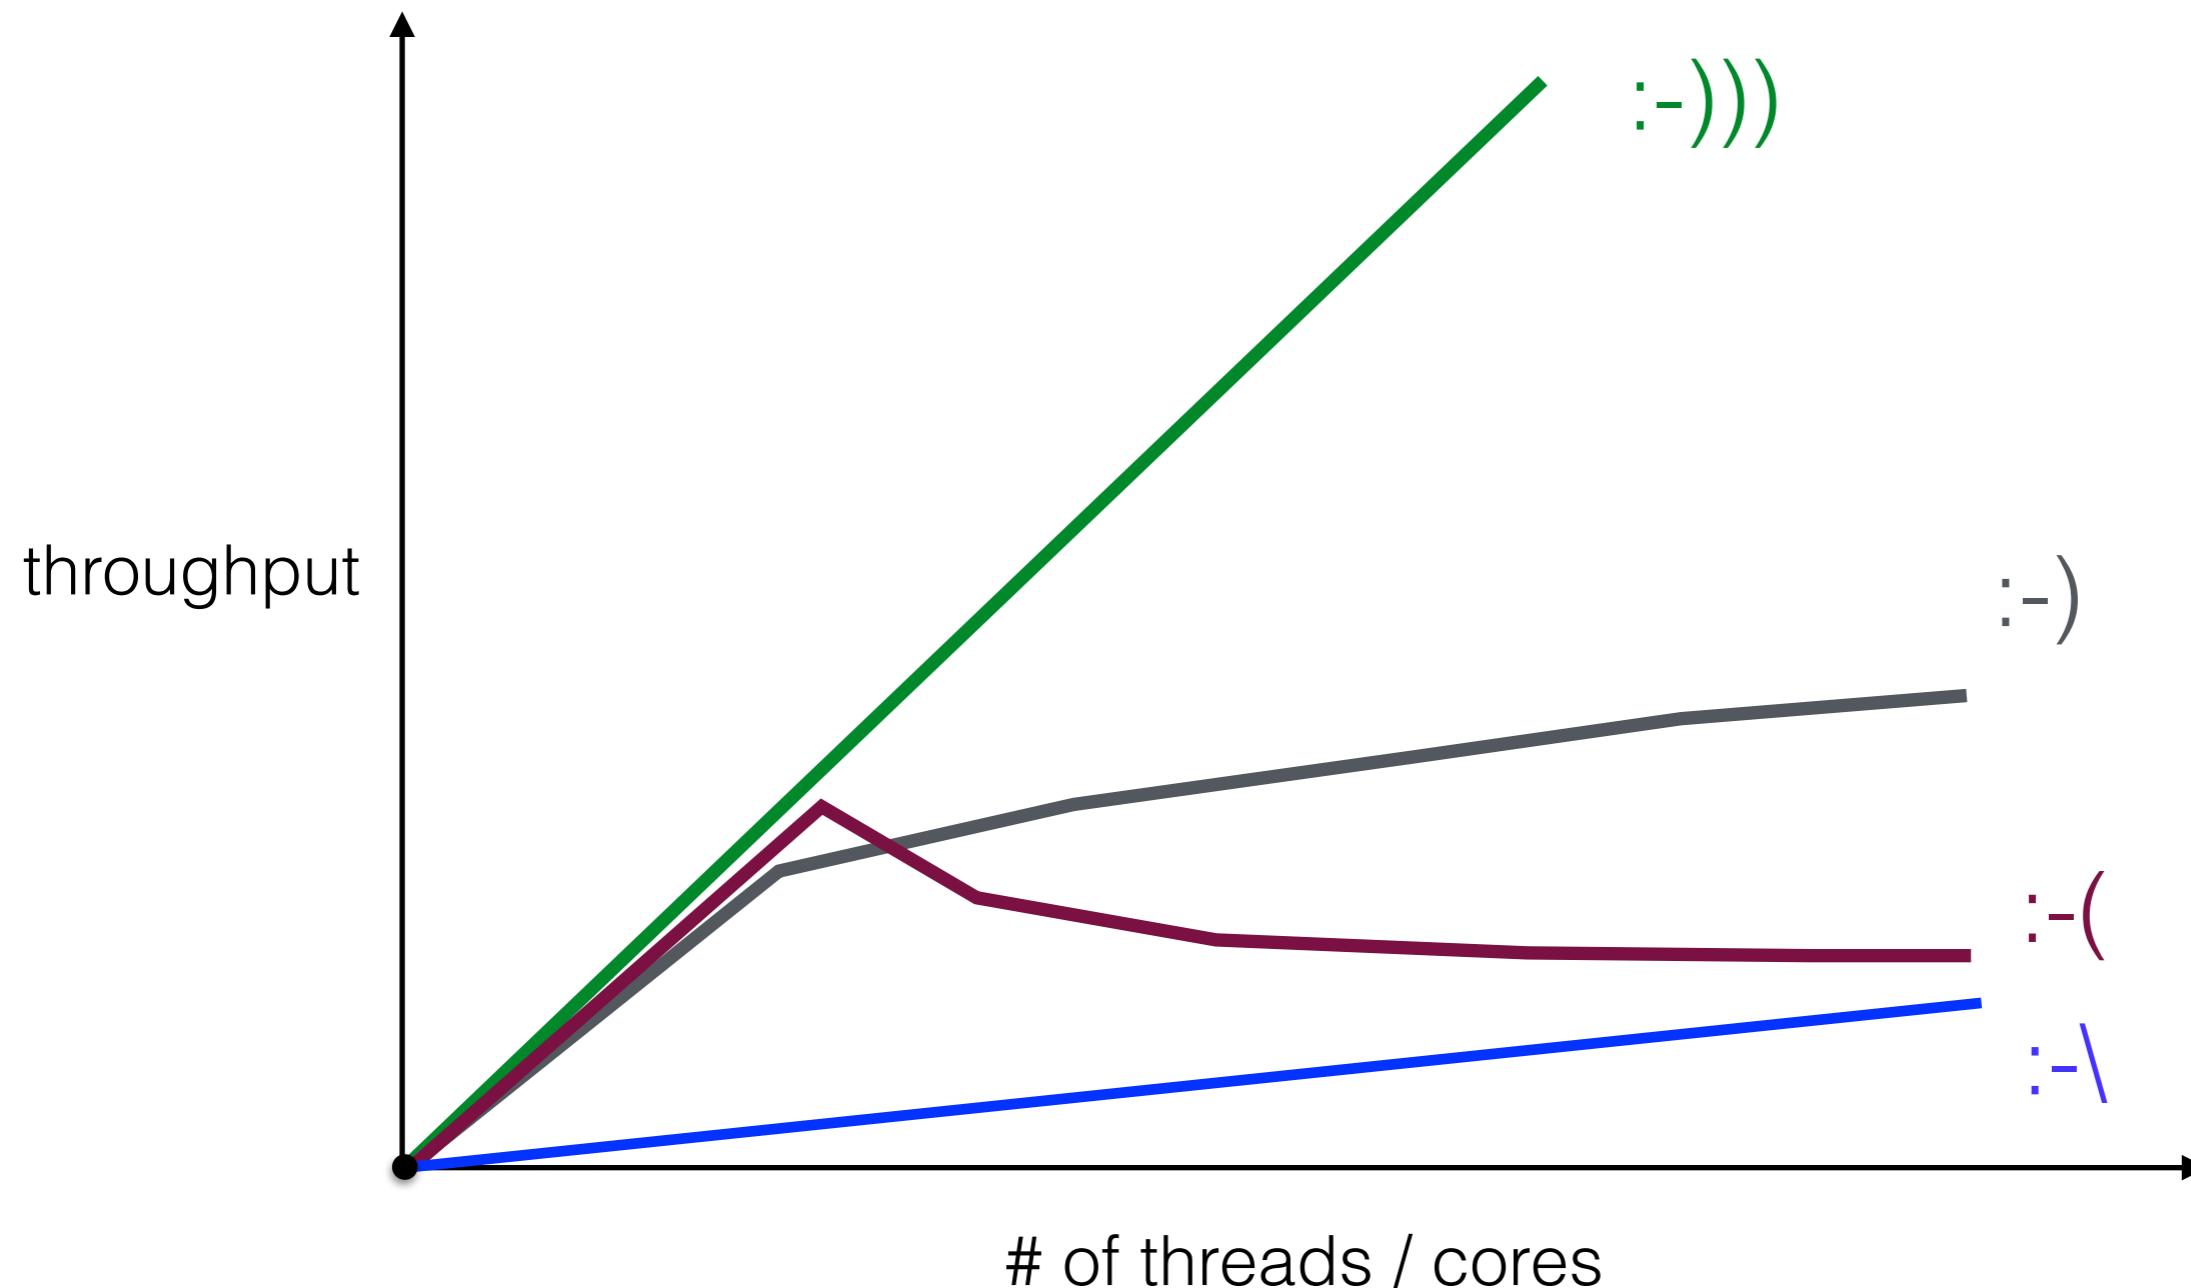
t1:	enq(2) ²	deq(1) ³
t2:	1 enq(1)	deq(2) ⁴

Sequential Consistency [Lamport'79]

there exists a sequential witness that preserves per-thread precedence (program order)

t1:	1 enq(1)	deq(2) ⁴
t2:	deq(1) ²	enq(2) ³

Performance and scalability



Relaxations allow trading
correctness
for
performance

Relaxations allow trading
correctness
for
performance

provide the potential
for better-performing
implementations

Relaxing the Semantics

Relaxing the Semantics

- Sequential specification = set of legal sequences
- Consistency condition = e.g. linearizability / sequential consistency

Relaxing the Semantics

Quantitative relaxations

Henzinger, Kirsch, Payer, Sezgin,S. POPL13

- Sequential specification = set of legal sequences
- Consistency condition = e.g. linearizability / sequential consistency

Relaxing the Semantics

not
“sequentially
correct”

Quantitative relaxations
Henzinger, Kirsch, Payer, Sezgin, S. POPL13

- Sequential specification = set of legal sequences
- Consistency condition = e.g. linearizability / sequential consistency

Relaxing the Semantics

not
“sequentially
correct”

Quantitative relaxations
Henzinger, Kirsch, Payer, Sezgin, S. POPL13

- Sequential specification = set of legal sequences
- Consistency condition = e.g. linearizability / sequential consistency

Local linearizability
in this talk

Relaxing the Semantics

not
“sequentially
correct”

Quantitative relaxations
Henzinger, Kirsch, Payer, Sezgin, S. POPL13

- Sequential specification = set of legal sequences
- Consistency condition = e.g. linearizability / sequential consistency

for queues only
(feel free to ask for more)

Local linearizability
in this talk

Relaxing the Semantics

not
“sequentially
correct”

Quantitative relaxations
Henzinger, Kirsch, Payer, Sezgin, S. POPL13

- Sequential specification = set of legal sequences
- Consistency condition = e.g. linearizability / sequential consistency

for queues only
(feel free to ask for more)

too weak

Local linearizability
in this talk

Local Linearizability

main idea

Local Linearizability

main idea

- Partition a history into a set of local histories
- Require linearizability per local history

Local Linearizability

main idea

Already present in some shared-memory consistency conditions
(not in our form of choice)

- Partition a history into a set of local histories
- Require linearizability per local history

Local Linearizability main idea

Already present in some shared-memory consistency conditions
(not in our form of choice)

- Partition a history into a set of local histories
- Require linearizability per local history

Local sequential consistency... is also possible

Local Linearizability main idea

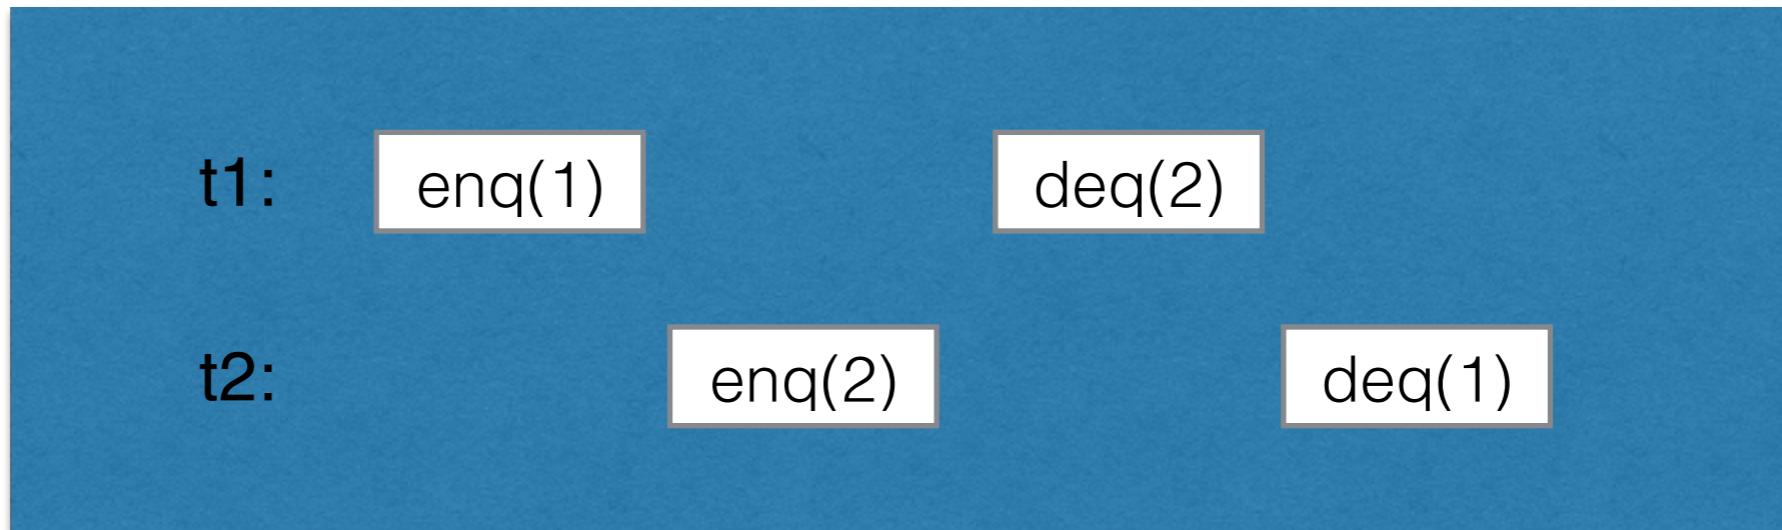
Already present in some shared-memory consistency conditions
(not in our form of choice)

- Partition a history into a set of local histories
- Require linearizability per local history

no global witness

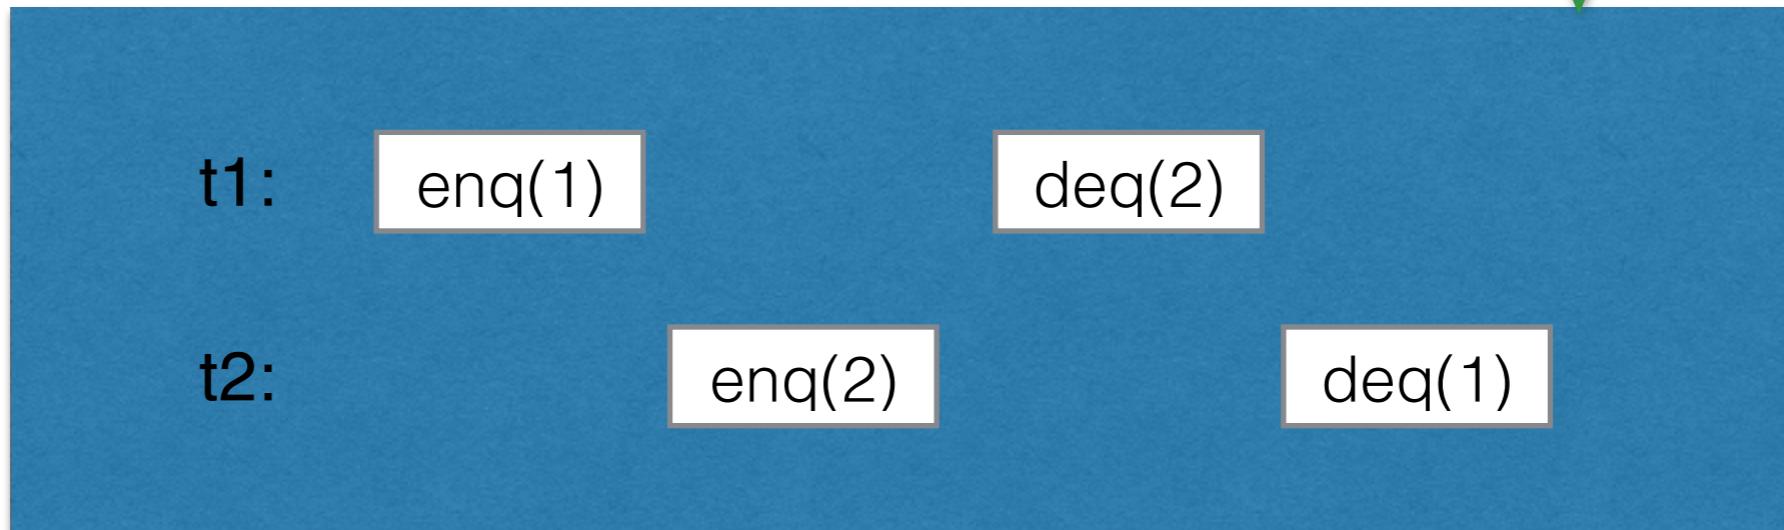
Local sequential consistency... is also possible

Local Linearizability (queue) example



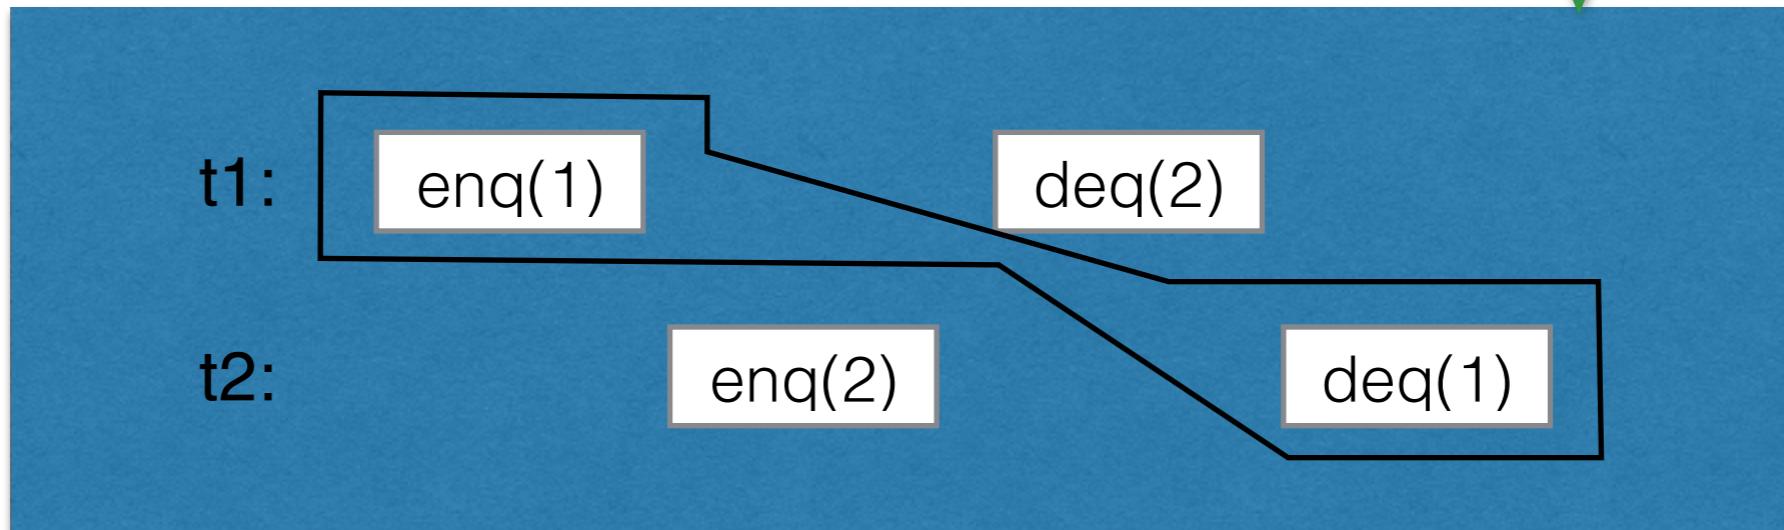
Local Linearizability (queue) example

(sequential) history
not linearizable



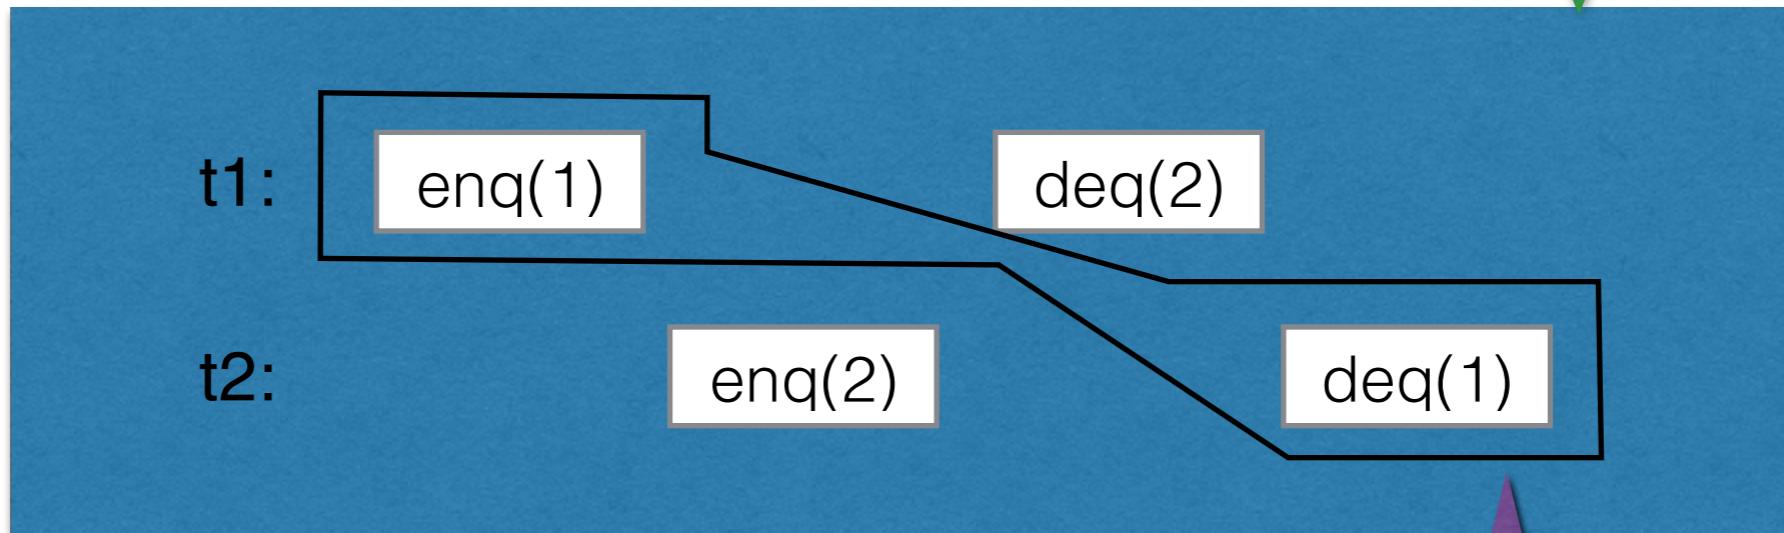
Local Linearizability (queue) example

(sequential) history
not linearizable



Local Linearizability (queue) example

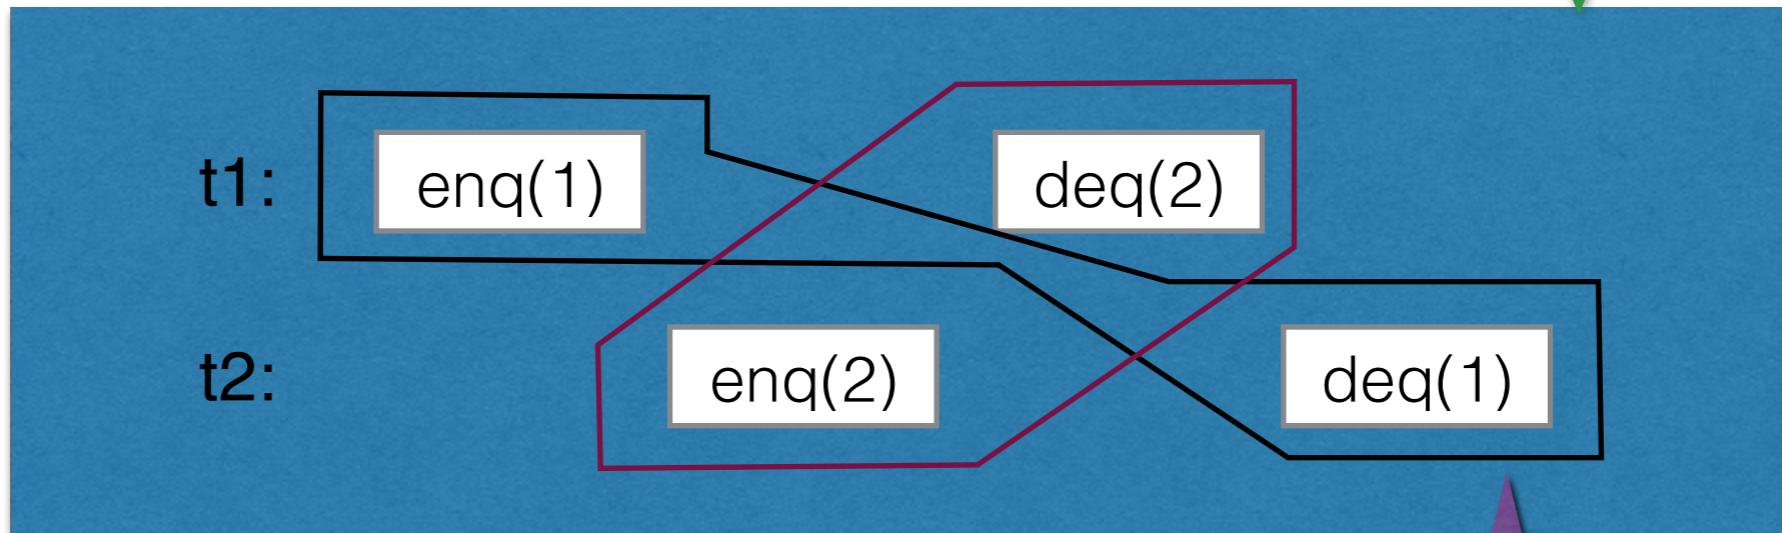
(sequential) history
not linearizable



t1-induced history,
linearizable

Local Linearizability (queue) example

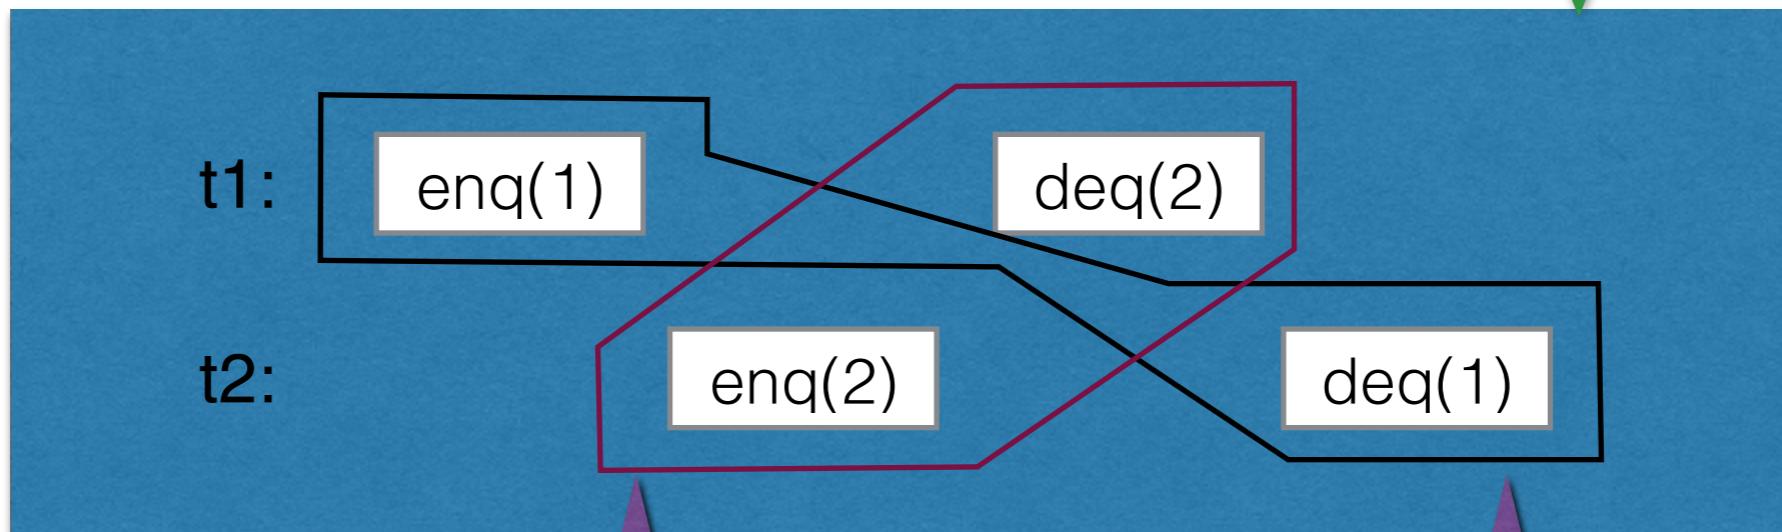
(sequential) history
not linearizable



t1-induced history,
linearizable

Local Linearizability (queue) example

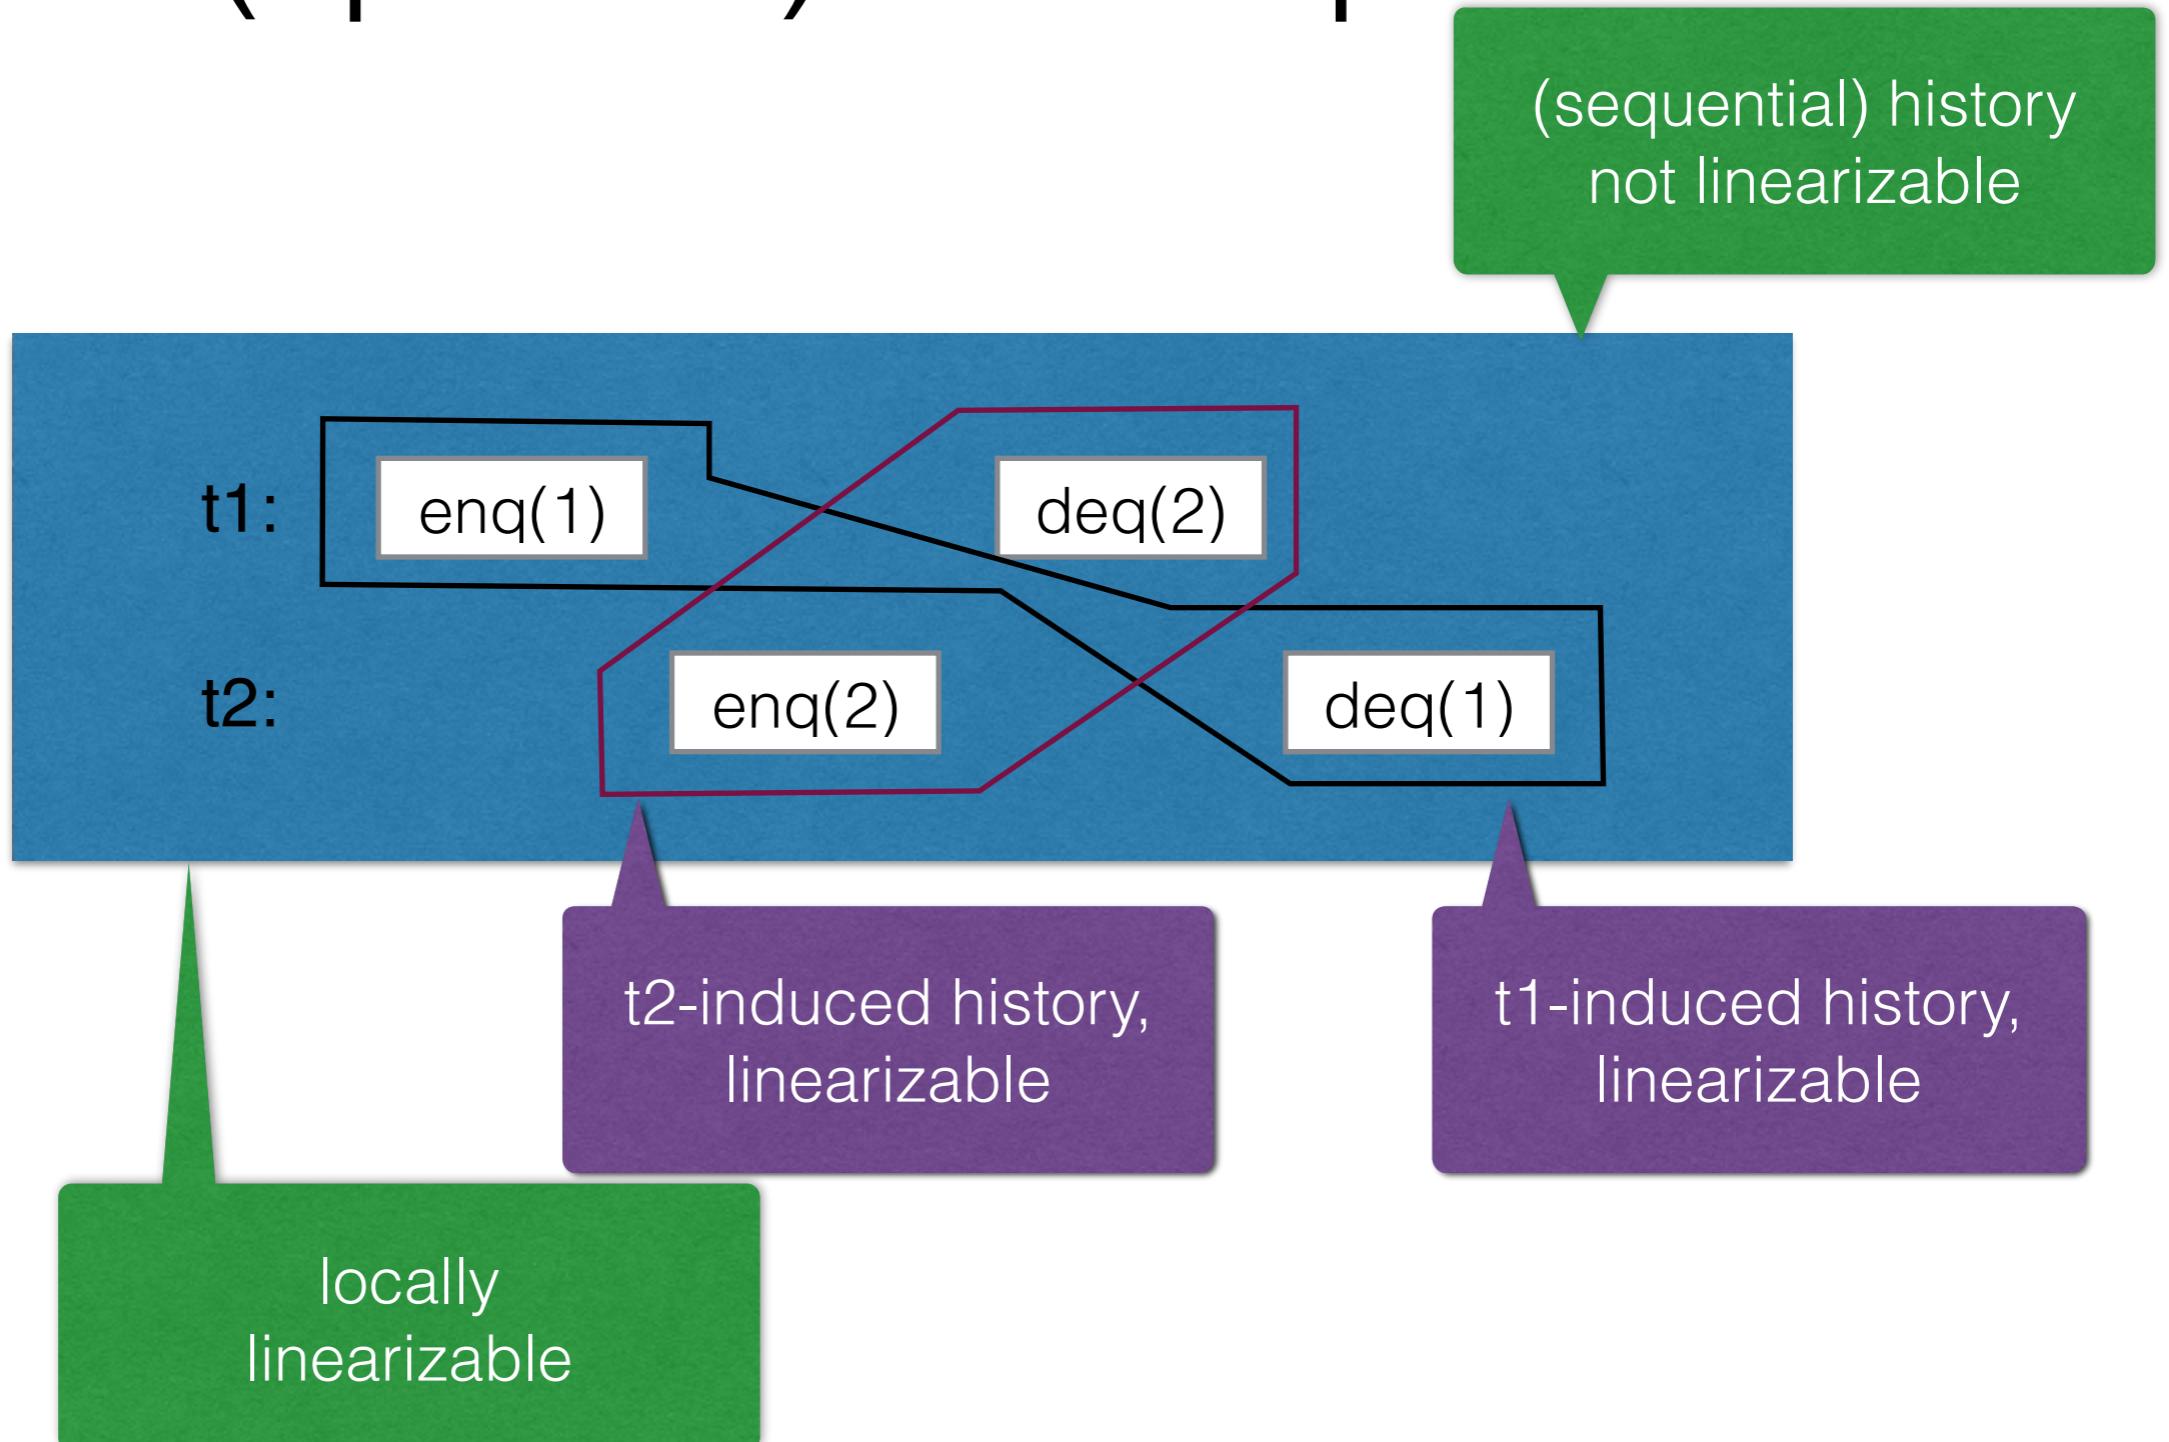
(sequential) history
not linearizable



t2-induced history,
linearizable

t1-induced history,
linearizable

Local Linearizability (queue) example



Local Linearizability (queue) definition

Local Linearizability (queue) definition

Queue signature $\Sigma = \{\text{enq}(x) \mid x \in V\} \cup \{\text{deq}(x) \mid x \in V\} \cup \{\text{deq}(\text{empty})\}$

Local Linearizability (queue) definition

Queue signature $\Sigma = \{\text{enq}(x) \mid x \in V\} \cup \{\text{deq}(x) \mid x \in V\} \cup \{\text{deq}(\text{empty})\}$

For a history \mathbf{h} with n threads, we put

$$\text{In}_{\mathbf{h}}(i) = \{\text{enq}(x)^i \in \mathbf{h} \mid x \in V\}$$

$$\text{Out}_{\mathbf{h}}(i) = \{\text{deq}(x)^i \in \mathbf{h} \mid \text{enq}(x)^i \in \text{In}_{\mathbf{h}}(i)\} \cup \{\text{deq}(\text{empty})\}$$

Local Linearizability (queue) definition

Queue signature $\Sigma = \{\text{enq}(x) \mid x \in V\} \cup \{\text{deq}(x) \mid x \in V\} \cup \{\text{deq}(\text{empty})\}$

For a history \mathbf{h} with n threads, we put

$$\text{In}_{\mathbf{h}}(i) = \{\text{enq}(x)^i \in \mathbf{h} \mid x \in V\}$$

in-methods of thread i
enqueues performed
by thread i

$$\text{Out}_{\mathbf{h}}(i) = \{\text{deq}(x)^i \in \mathbf{h} \mid \text{enq}(x)^i \in \text{In}_{\mathbf{h}}(i)\} \cup \{\text{deq}(\text{empty})\}$$

Local Linearizability (queue) definition

Queue signature $\Sigma = \{\text{enq}(x) \mid x \in V\} \cup \{\text{deq}(x) \mid x \in V\} \cup \{\text{deq}(\text{empty})\}$

For a history \mathbf{h} with n threads, we put

$$\text{In}_{\mathbf{h}}(i) = \{\text{enq}(x)^i \in \mathbf{h} \mid x \in V\}$$

in-methods of thread i
enqueues performed
by thread i

$$\text{Out}_{\mathbf{h}}(i) = \{\text{deq}(x)^j \in \mathbf{h} \mid \text{enq}(x)^i \in \text{In}_{\mathbf{h}}(i)\} \cup \{\text{deq}(\text{empty})\}$$

out-methods of thread i
dequeues
(performed by any thread)
corresponding to enqueues that
are in-methods

Local Linearizability (queue) definition

Queue signature $\Sigma = \{\text{enq}(x) \mid x \in V\} \cup \{\text{deq}(x) \mid x \in V\} \cup \{\text{deq}(\text{empty})\}$

For a history \mathbf{h} with n threads, we put

$$\text{In}_{\mathbf{h}}(i) = \{\text{enq}(x)^i \in \mathbf{h} \mid x \in V\}$$

in-methods of thread i
enqueues performed
by thread i

$$\text{Out}_{\mathbf{h}}(i) = \{\text{deq}(x)^j \in \mathbf{h} \mid \text{enq}(x)^i \in \text{In}_{\mathbf{h}}(i)\} \cup \{\text{deq}(\text{empty})\}$$

out-methods of thread i
dequeues
(performed by any thread)
corresponding to enqueues that
are in-methods

\mathbf{h} is locally linearizable iff every thread-induced history

$$\mathbf{h}_i = \mathbf{h} \mid (\text{In}_{\mathbf{h}}(i) \cup \text{Out}_{\mathbf{h}}(i))$$

is linearizable.

Generalizations of Local Linearizability

Generalizations of Local Linearizability

Signature Σ

Generalizations of Local Linearizability

Signature Σ

For a history \mathbf{h} with n threads, choose

$\text{In}_{\mathbf{h}}(i)$

$\text{Out}_{\mathbf{h}}(i)$

Generalizations of Local Linearizability

Signature Σ

For a history \mathbf{h} with n threads, choose

$In_{\mathbf{h}}(i)$

$Out_{\mathbf{h}}(i)$

in-methods of thread i ,
methods that go in \mathbf{h}_i

Generalizations of Local Linearizability

Signature Σ

For a history \mathbf{h} with n threads, choose

$In_{\mathbf{h}}(i)$

$Out_{\mathbf{h}}(i)$

in-methods of thread i ,
methods that go in \mathbf{h}_i

out-methods of thread i ,
dependent methods
on the methods in $In_{\mathbf{h}}(i)$
(performed by any thread)

Generalizations of Local Linearizability

Signature Σ

For a history \mathbf{h} with n threads, choose

$\text{In}_{\mathbf{h}}(i)$

$\text{Out}_{\mathbf{h}}(i)$

in-methods of thread i ,
methods that go in \mathbf{h}_i

out-methods of thread i ,
dependent methods
on the methods in $\text{In}_{\mathbf{h}}(i)$
(performed by any thread)

\mathbf{h} is locally linearizable iff every thread-induced history

$$\mathbf{h}_i = \mathbf{h} \mid (\text{In}_{\mathbf{h}}(i) \cup \text{Out}_{\mathbf{h}}(i))$$

is linearizable.

Generalizations of Local Linearizability

Signature Σ

For a history \mathbf{h} with n threads, choose

$\text{In}_{\mathbf{h}}(i)$

$\text{Out}_{\mathbf{h}}(i)$

by increasing the
in-methods,
LL gradually moves to
linearizability

in-methods of thread i ,
methods that go in \mathbf{h}_i

out-methods of thread i ,
dependent methods
on the methods in $\text{In}_{\mathbf{h}}(i)$
(performed by any thread)

\mathbf{h} is locally linearizable iff every thread-induced history

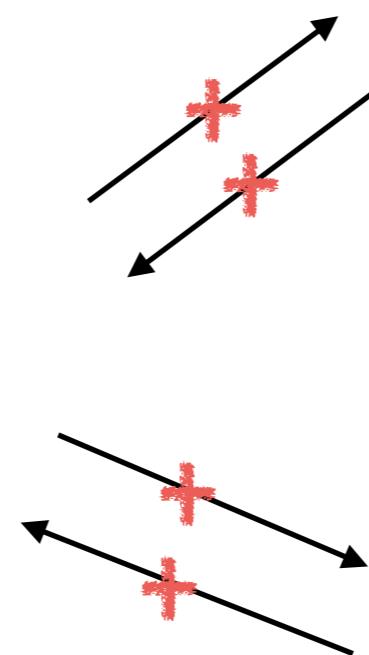
$$\mathbf{h}_i = \mathbf{h} \mid (\text{In}_{\mathbf{h}}(i) \cup \text{Out}_{\mathbf{h}}(i))$$

is linearizable.

Where do we stand?

In general

Local Linearizability



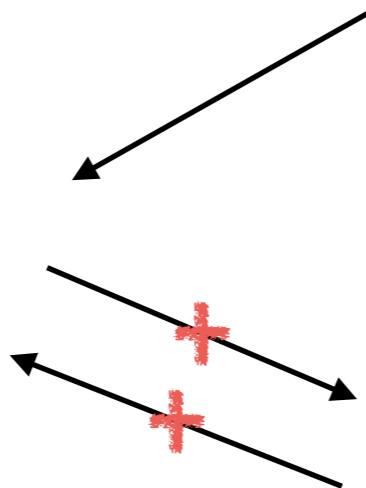
Linearizability

Sequential Consistency

Where do we stand?

For queues (and all pool-like data structures)

Local Linearizability



Linearizability

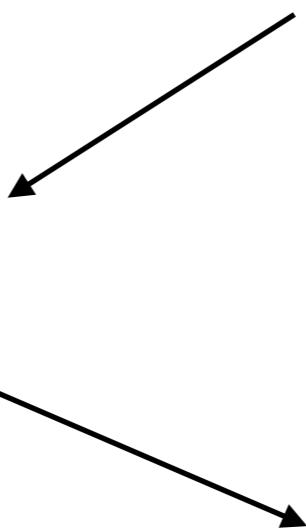


Sequential Consistency

Where do we stand?

C: For queues

Local Linearizability
& Pool-seq.cons.



Linearizability

Sequential Consistency

Properties

Properties

Local linearizability is compositional

Properties

Local linearizability is compositional

like linearizability
unlike sequential consistency

Properties

Local linearizability is compositional

like linearizability
unlike sequential consistency

h (over multiple objects) is locally linearizable
iff

each per-object subhistory of **h** is locally linearizable

Properties

Local linearizability is compositional

like linearizability
unlike sequential consistency

h (over multiple objects) is locally linearizable
iff

each per-object subhistory of **h** is locally linearizable

Local linearizability is modular /
“decompositional”

Properties

Local linearizability is compositional

like linearizability
unlike sequential consistency

h (over multiple objects) is locally linearizable
iff

each per-object subhistory of **h** is locally linearizable

Local linearizability is modular /
“decompositional”

uses decomposition into smaller
histories, by definition

Properties

Local linearizability is compositional

like linearizability
unlike sequential consistency

h (over multiple objects) is locally linearizable
iff

each per-object subhistory of **h** is locally linearizable

Local linearizability is modular /
“decompositional”

uses decomposition into smaller histories, by definition

may allow for modular verification

Verification (queue)

Queue sequential specification (axiomatic)

s is a legal queue sequence
iff

1. **s** is a legal pool sequence, and
2. $\text{enq}(x) <_s \text{enq}(y) \wedge \text{deq}(y) \in \mathbf{s} \Rightarrow \text{deq}(x) \in \mathbf{s} \wedge \text{deq}(x) <_s \text{deq}(y)$

Verification (queue)

Queue sequential specification (axiomatic)

s is a legal queue sequence
iff

1. **s** is a legal pool sequence, and
2. $\text{enq}(x) <_s \text{enq}(y) \wedge \text{deq}(y) \in \mathbf{s} \Rightarrow \text{deq}(x) \in \mathbf{s} \wedge \text{deq}(x) <_s \text{deq}(y)$

Queue linearizability (axiomatic)

h is queue linearizable
iff

1. **h** is pool linearizable, and
2. $\text{enq}(x) <_h \text{enq}(y) \wedge \text{deq}(y) \in \mathbf{h} \Rightarrow \text{deq}(x) \in \mathbf{h} \wedge \text{deq}(y) <_h \text{deq}(x)$

Verification (queue)

Queue sequential specification (axiomatic)

s is a legal queue sequence
iff

1. **s** is a legal pool sequence, and
2. $\text{enq}(x) <_s \text{enq}(y) \wedge \text{deq}(y) \in \mathbf{s} \Rightarrow \text{deq}(x) \in \mathbf{s} \wedge \text{deq}(x) <_s \text{deq}(y)$

Queue linearizability (axiomatic)

h is queue linearizable
iff

1. **h** is pool linearizable, and
2. $\text{enq}(x) <_h \text{enq}(y) \wedge \text{deq}(y) \in \mathbf{h} \Rightarrow \text{deq}(x) \in \mathbf{h} \wedge \text{deq}(y) <_h \text{deq}(x)$

precedence order

Verification (queue)

Queue sequential specification (axiomatic)

s is a legal queue sequence
iff

1. **s** is a legal pool sequence, and
2. $\text{enq}(x) <_s \text{enq}(y) \wedge \text{deq}(y) \in \mathbf{s} \Rightarrow \text{deq}(x) \in \mathbf{s} \wedge \text{deq}(x) <_s \text{deq}(y)$

Verification (queue)

Queue sequential specification (axiomatic)

s is a legal queue sequence
iff

1. **s** is a legal pool sequence, and
2. $\text{enq}(x) <_s \text{enq}(y) \wedge \text{deq}(y) \in \mathbf{s} \Rightarrow \text{deq}(x) \in \mathbf{s} \wedge \text{deq}(x) <_s \text{deq}(y)$

Queue local linearizability (axiomatic)

h is queue locally linearizable
iff

1. **h** is pool locally linearizable, and
2. $\text{enq}(x) <_{\mathbf{h}} \text{enq}(y) \wedge \text{deq}(y) \in \mathbf{h} \Rightarrow \text{deq}(x) \in \mathbf{h} \wedge \text{deq}(y) <_{\mathbf{h}} \text{deq}(x)$

Verification (queue)

Queue sequential specification (axiomatic)

s is a legal queue sequence
iff

1. **s** is a legal pool sequence, and
2. $\text{enq}(x) <_s \text{enq}(y) \wedge \text{deq}(y) \in \mathbf{s} \Rightarrow \text{deq}(x) \in \mathbf{s} \wedge \text{deq}(x) <_s \text{deq}(y)$

Queue local linearizability (axiomatic)

h is queue locally linearizable
iff

1. **h** is pool locally linearizable, and
2. $\text{enq}(x) <_h \text{enq}(y) \wedge \text{deq}(y) \in \mathbf{h} \Rightarrow \text{deq}(x) \in \mathbf{h} \wedge \text{deq}(y) <_h \text{deq}(x)$

thread-local
precedence order

Generic Implementations

Generic Implementations

Your favorite linearizable data structure implementation

Generic Implementations

Your favorite linearizable data structure implementation

\emptyset

Generic Implementations

Your favorite linearizable data structure implementation

Φ

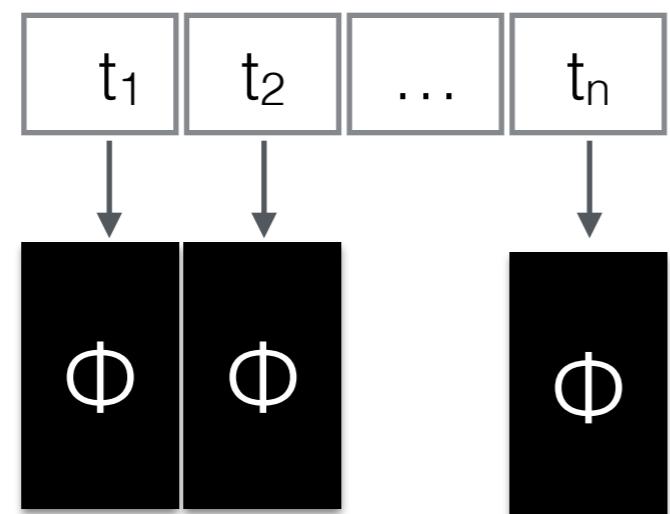
turns into a locally linearizable implementation by:

Generic Implementations

Your favorite linearizable data structure implementation



turns into a locally linearizable implementation by:

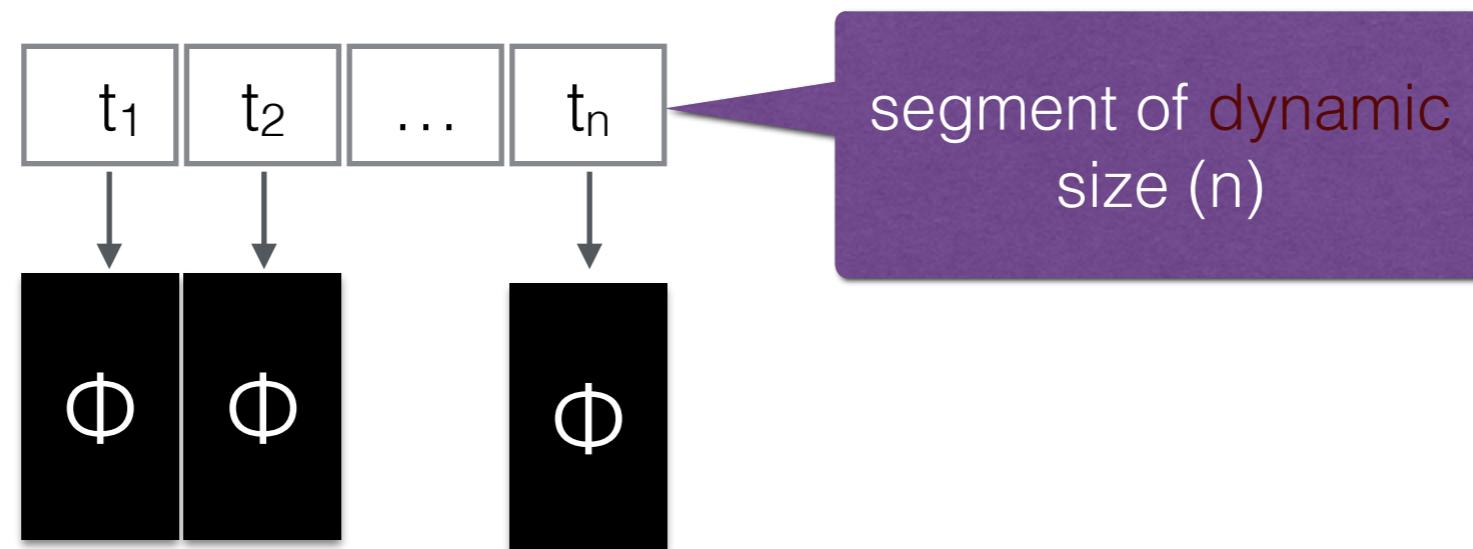


Generic Implementations

Your favorite linearizable data structure implementation



turns into a locally linearizable implementation by:

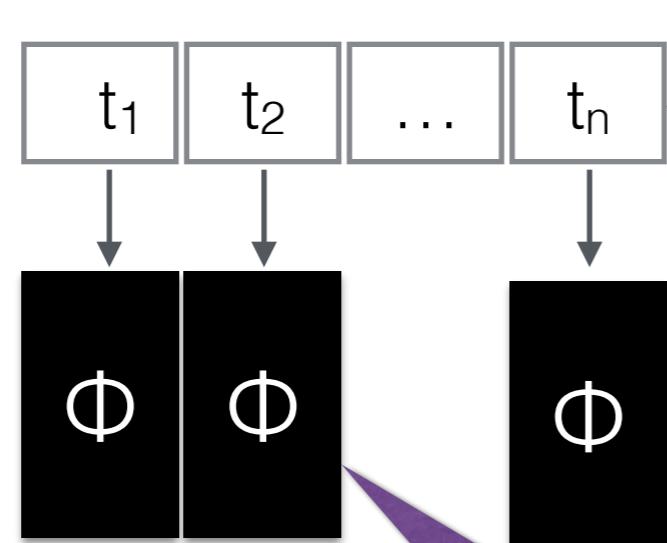


Generic Implementations

Your favorite linearizable data structure implementation



turns into a locally linearizable implementation by:



segment of dynamic
size (n)

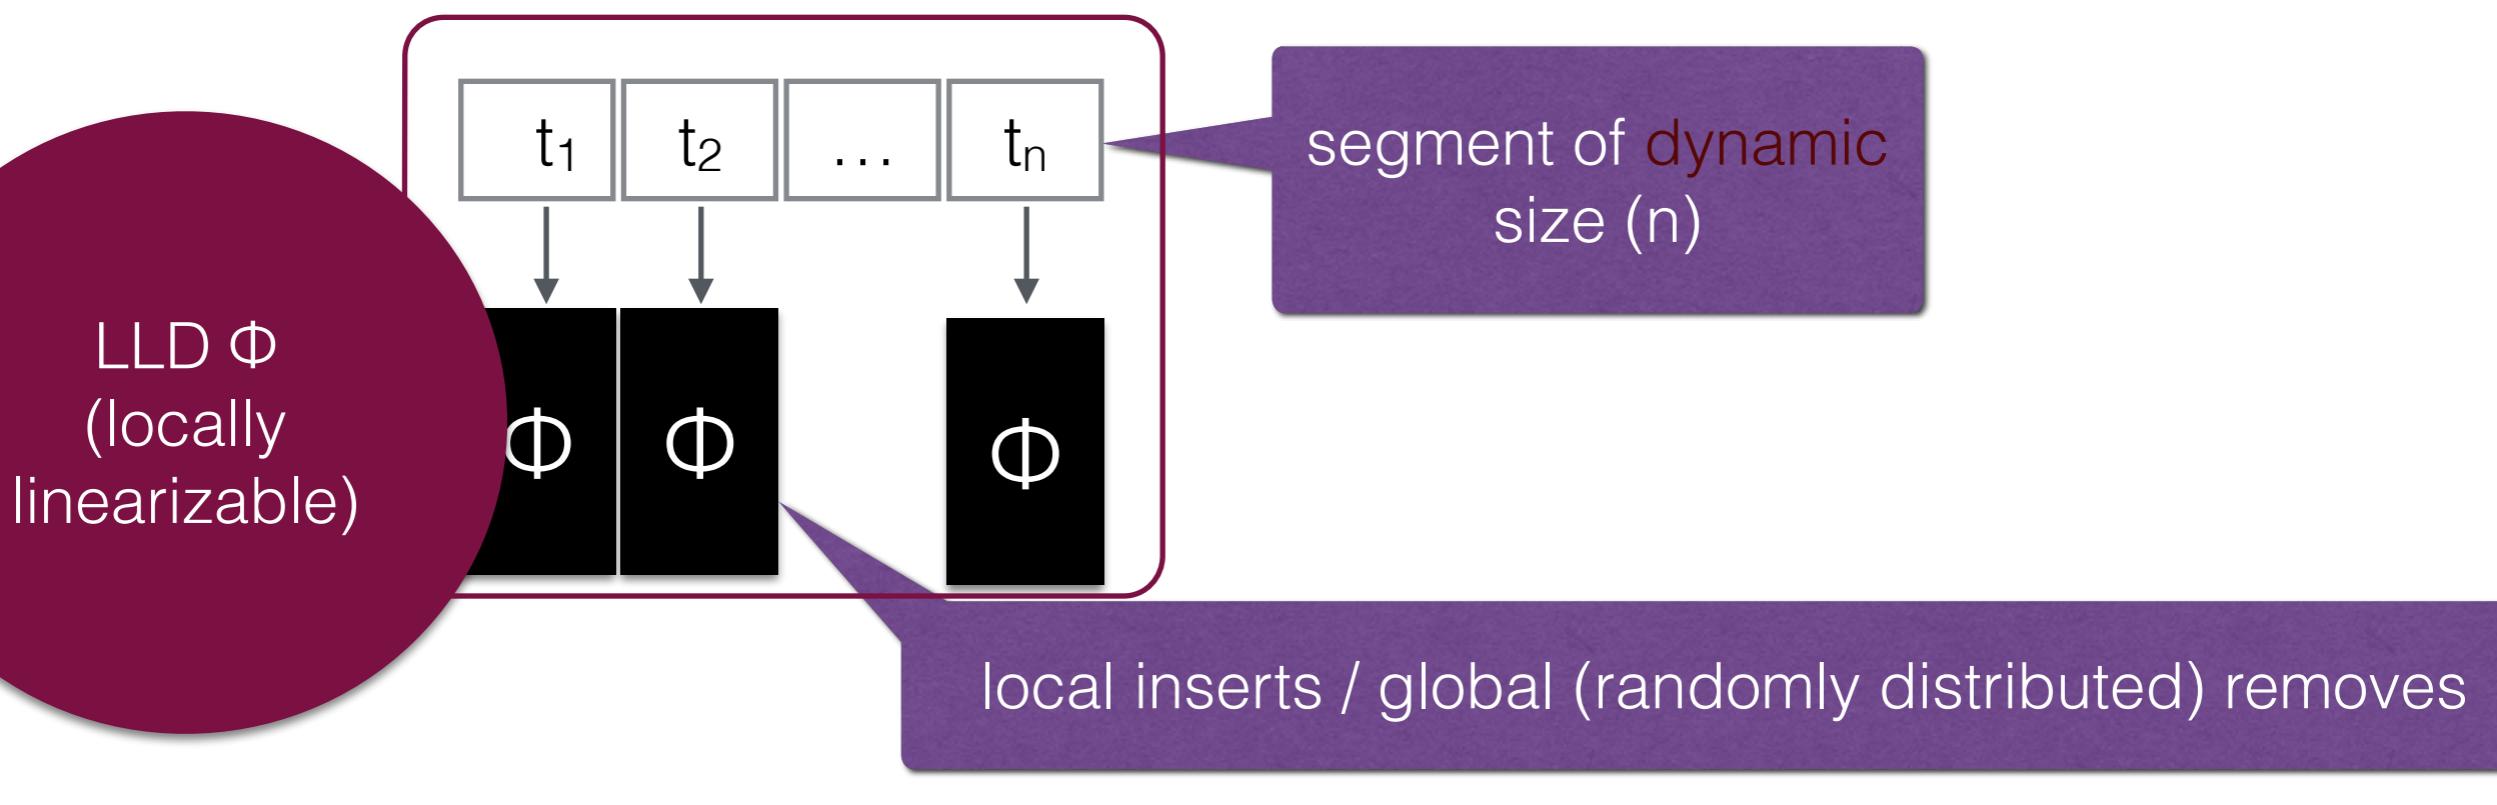
local inserts / global (randomly distributed) removes

Generic Implementations

Your favorite linearizable data structure implementation



turns into a locally linearizable implementation by:



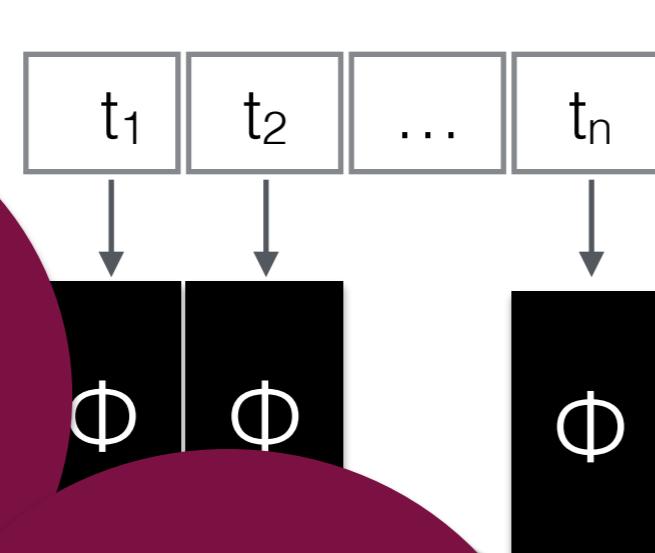
Generic Implementations

Your favorite linearizable data structure implementation

Φ

turns into a locally linearizable implementation by:

LLD Φ
(locally
linearizable)

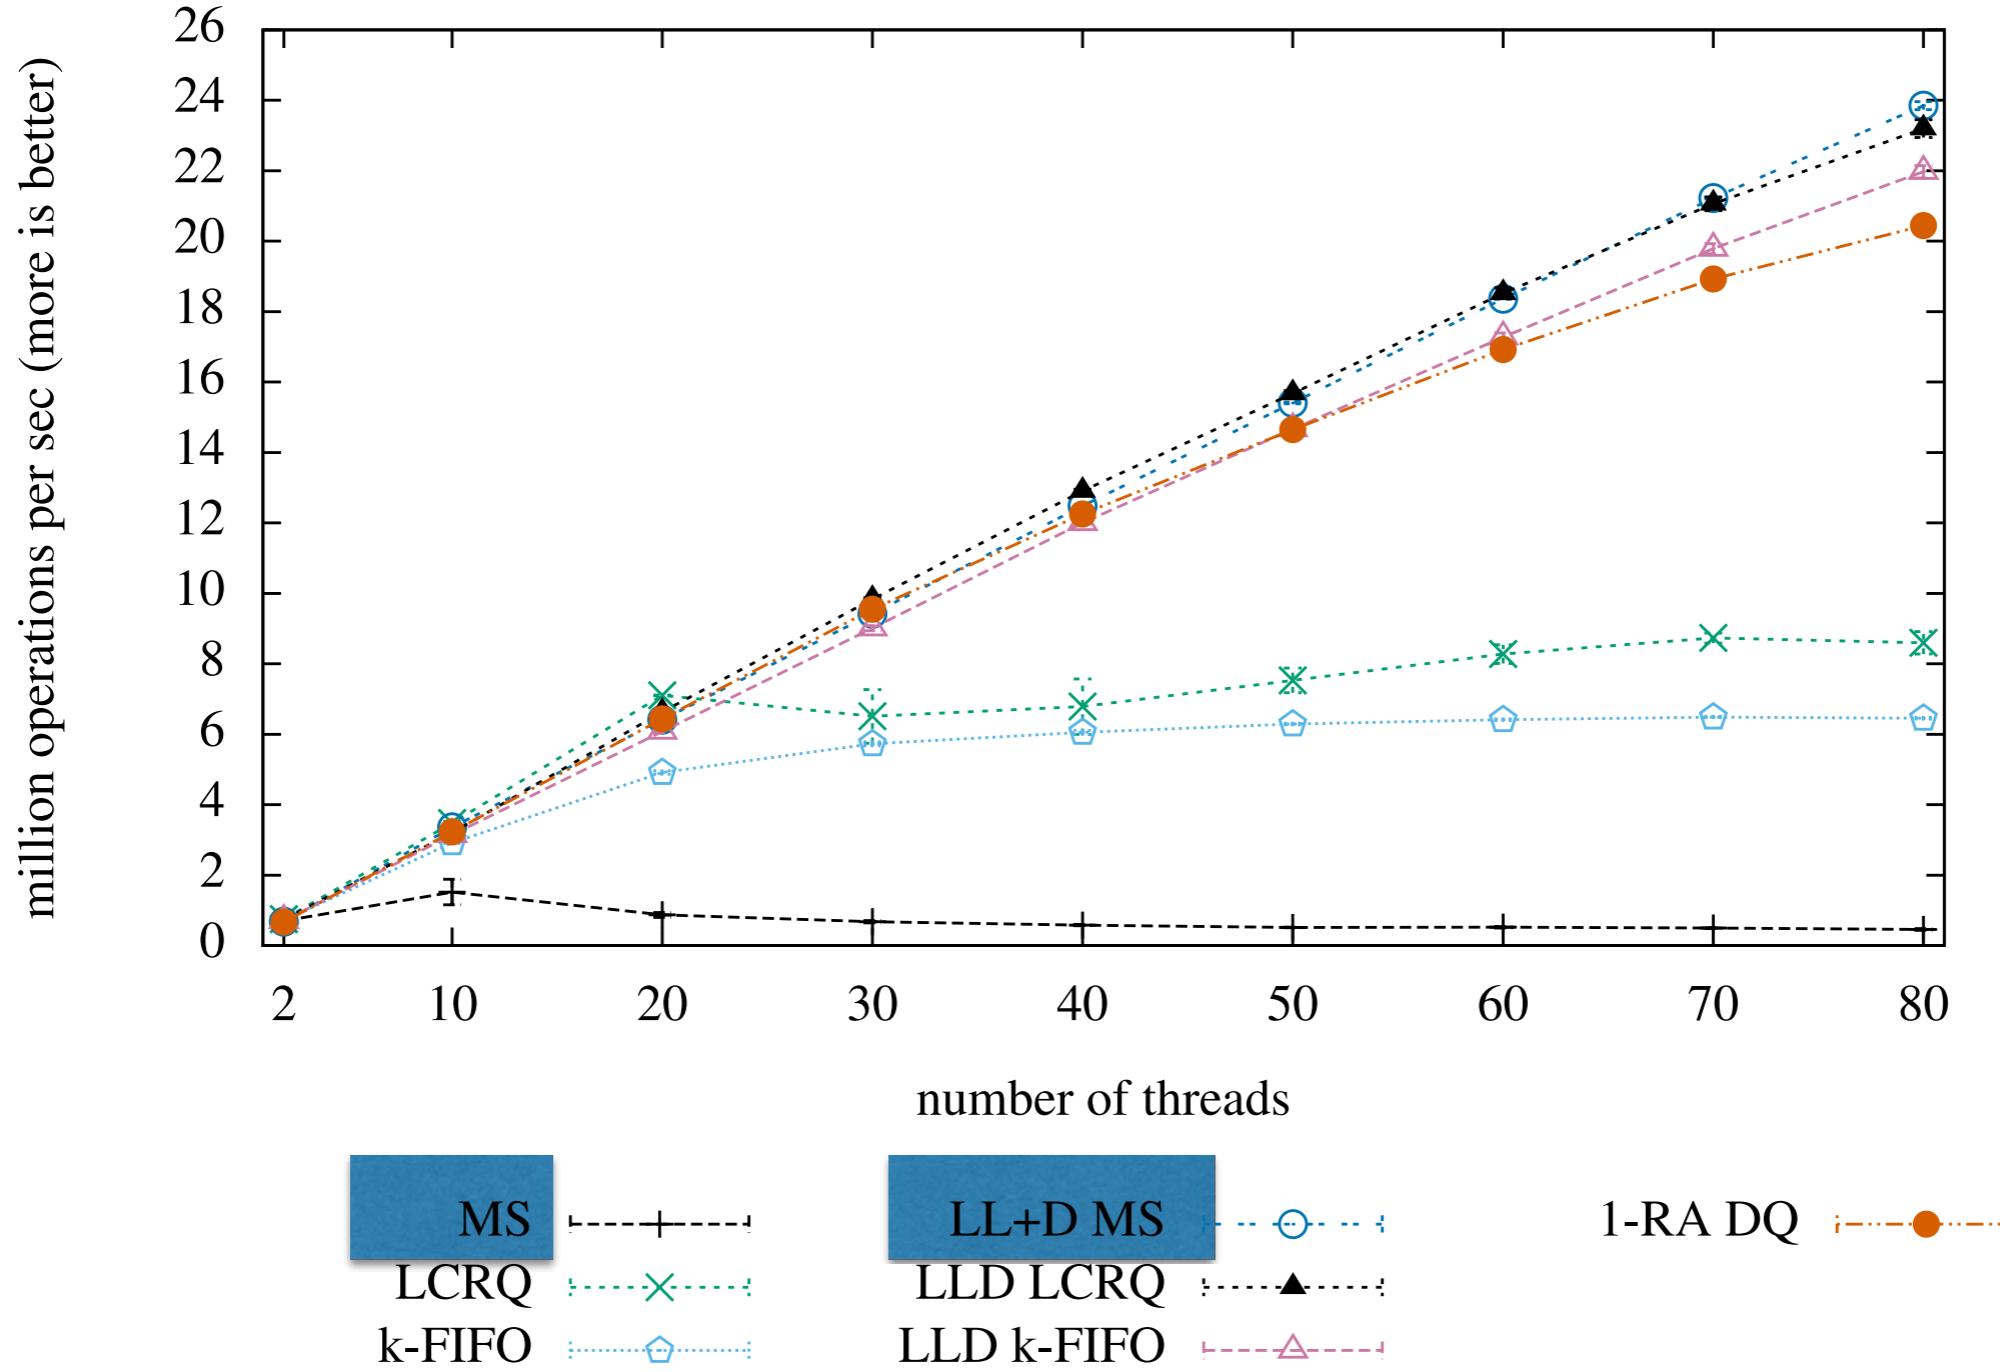


segment of dynamic
size (n)

LL+D Φ
(also pool
linearizable)

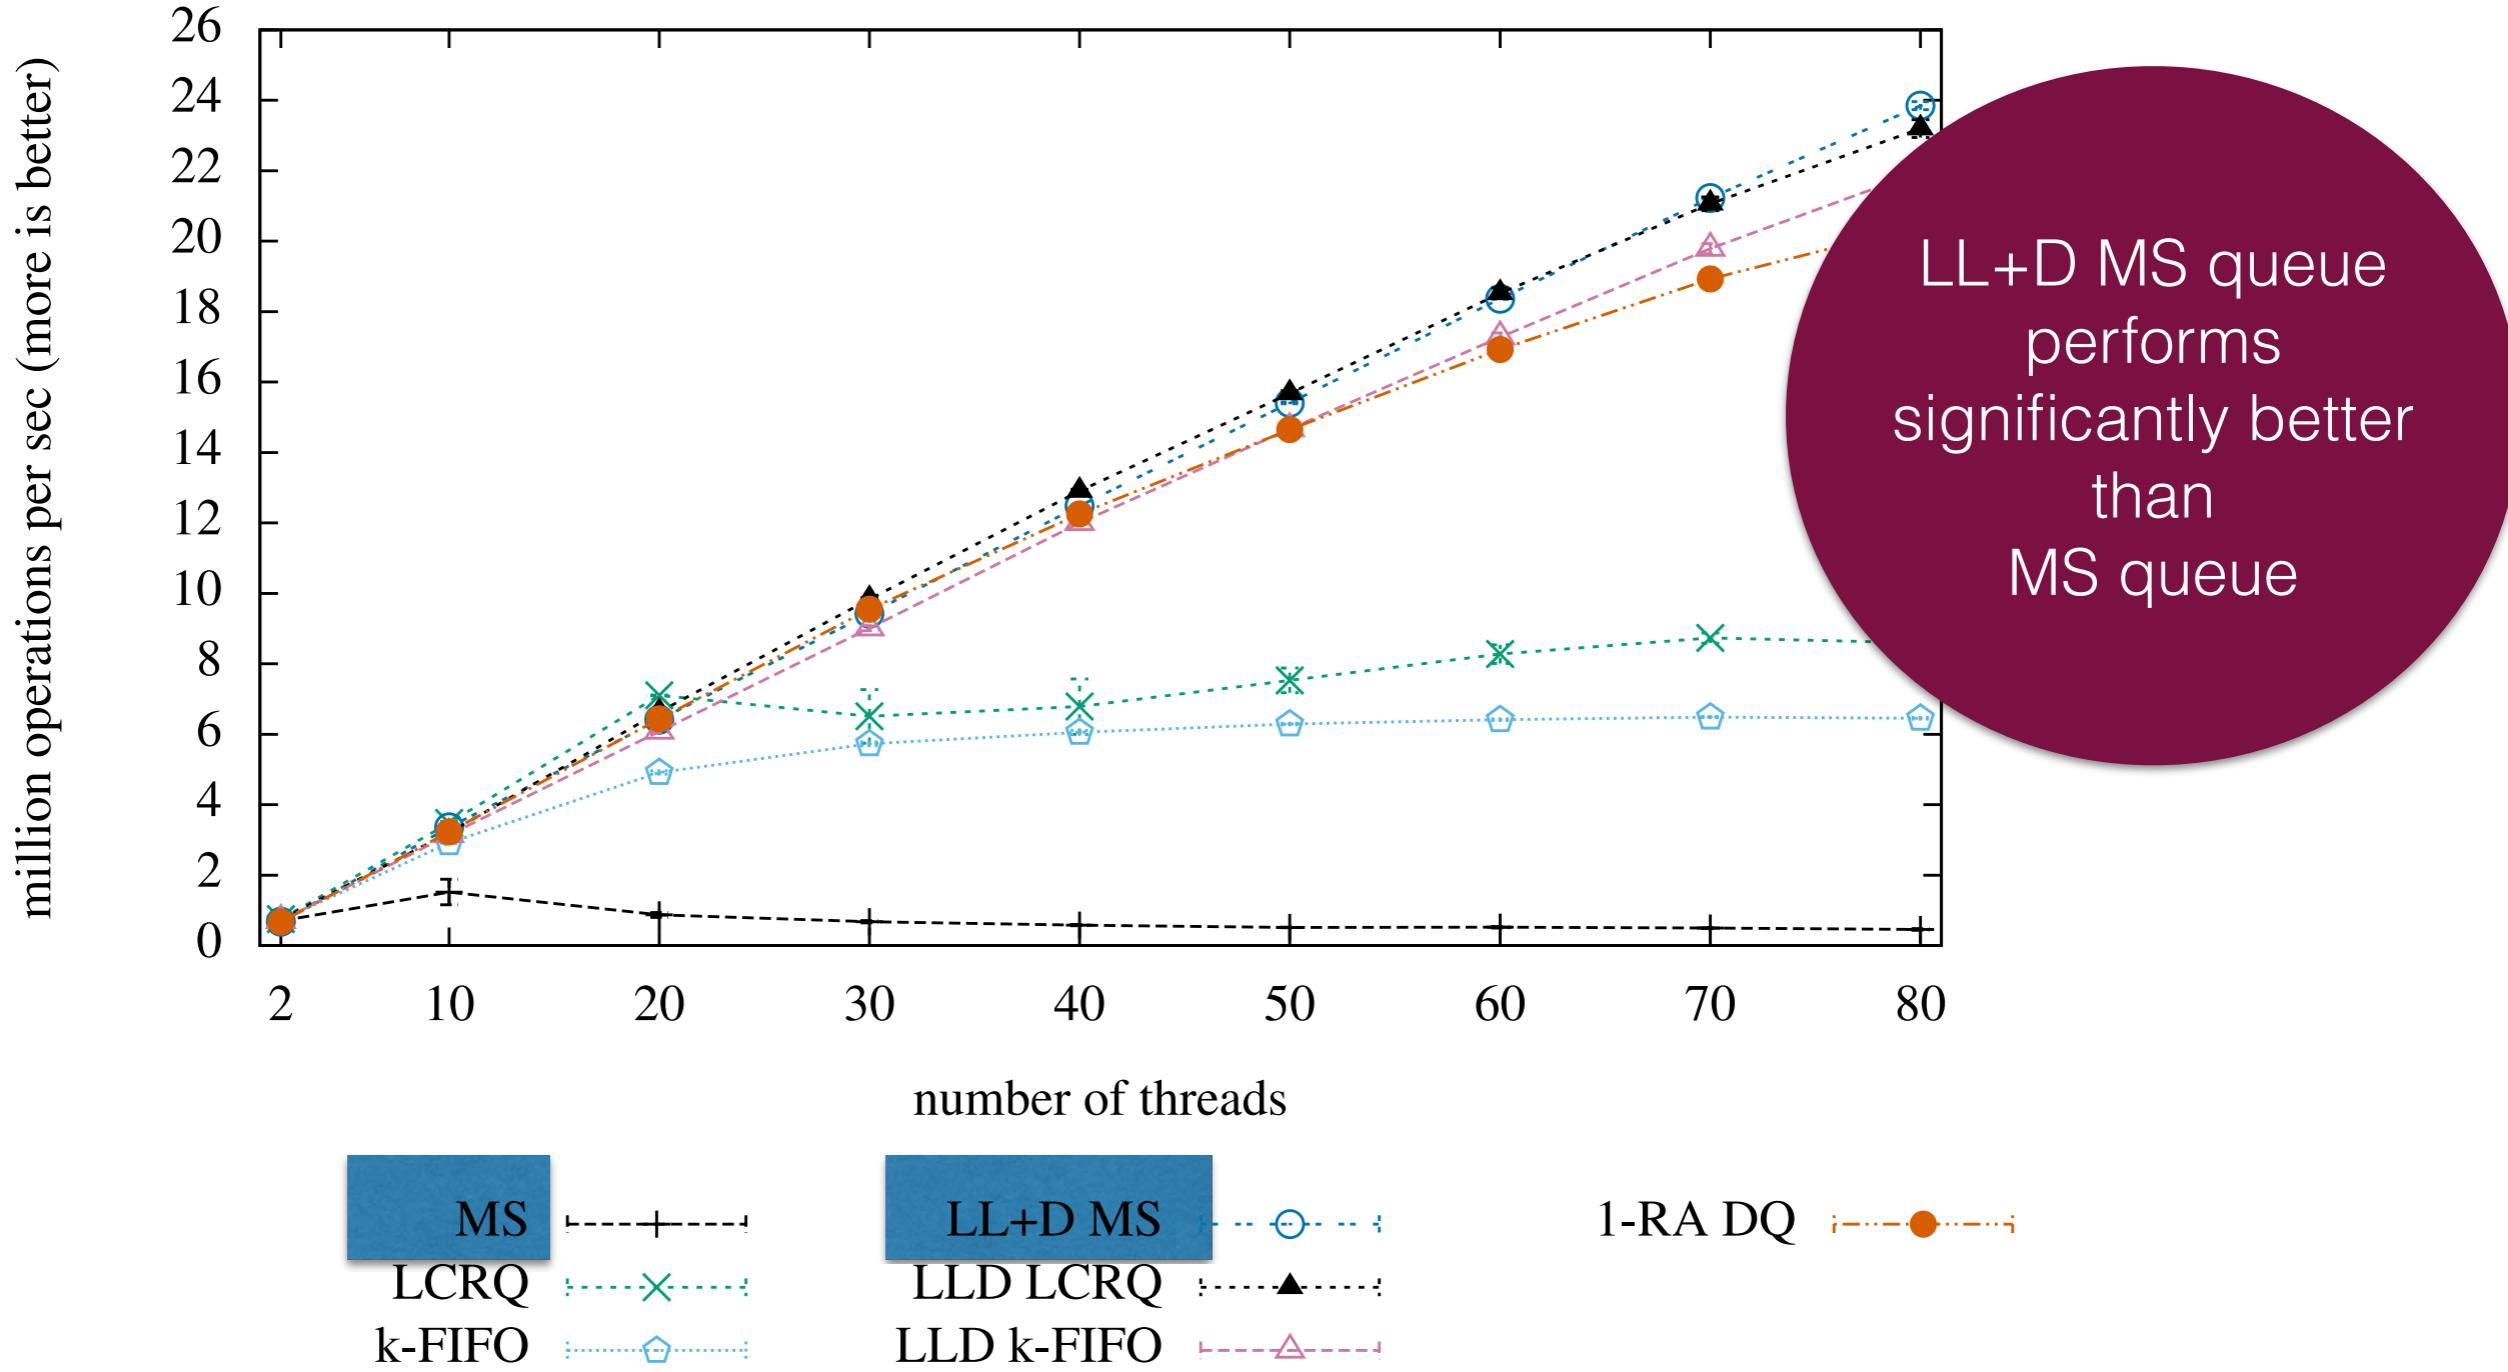
local inserts / global (randomly distributed) removes

Performance



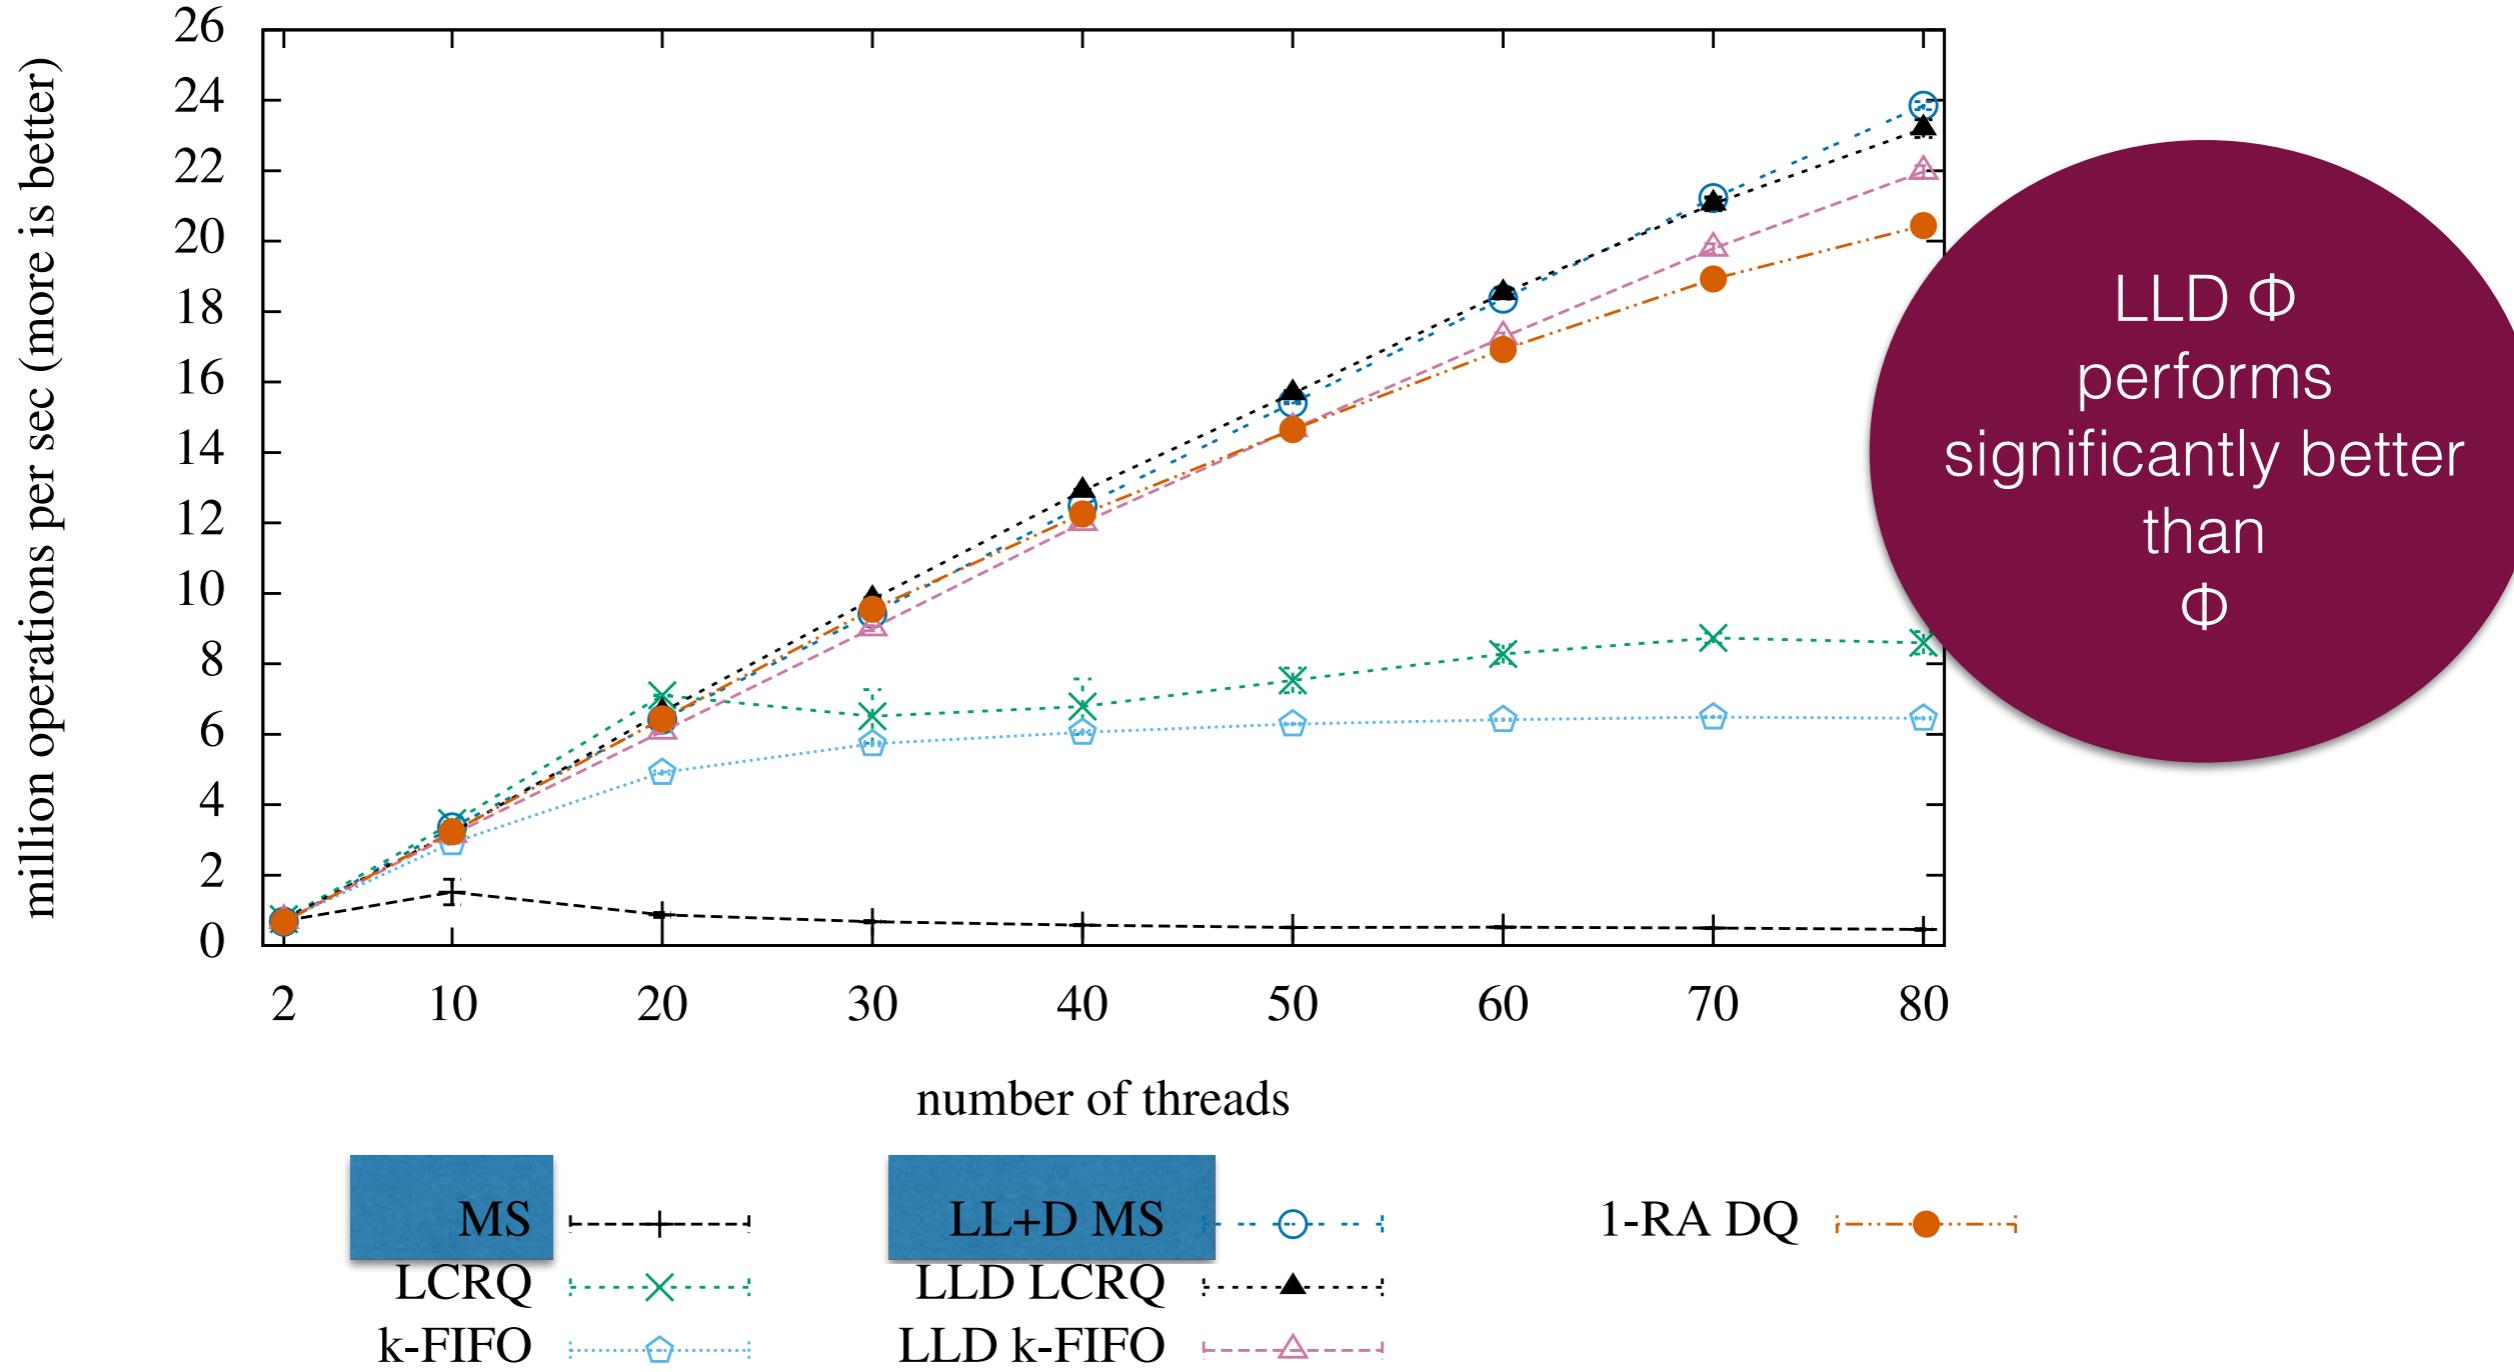
(a) Queues, LL queues, and “queue-like” pools

Performance



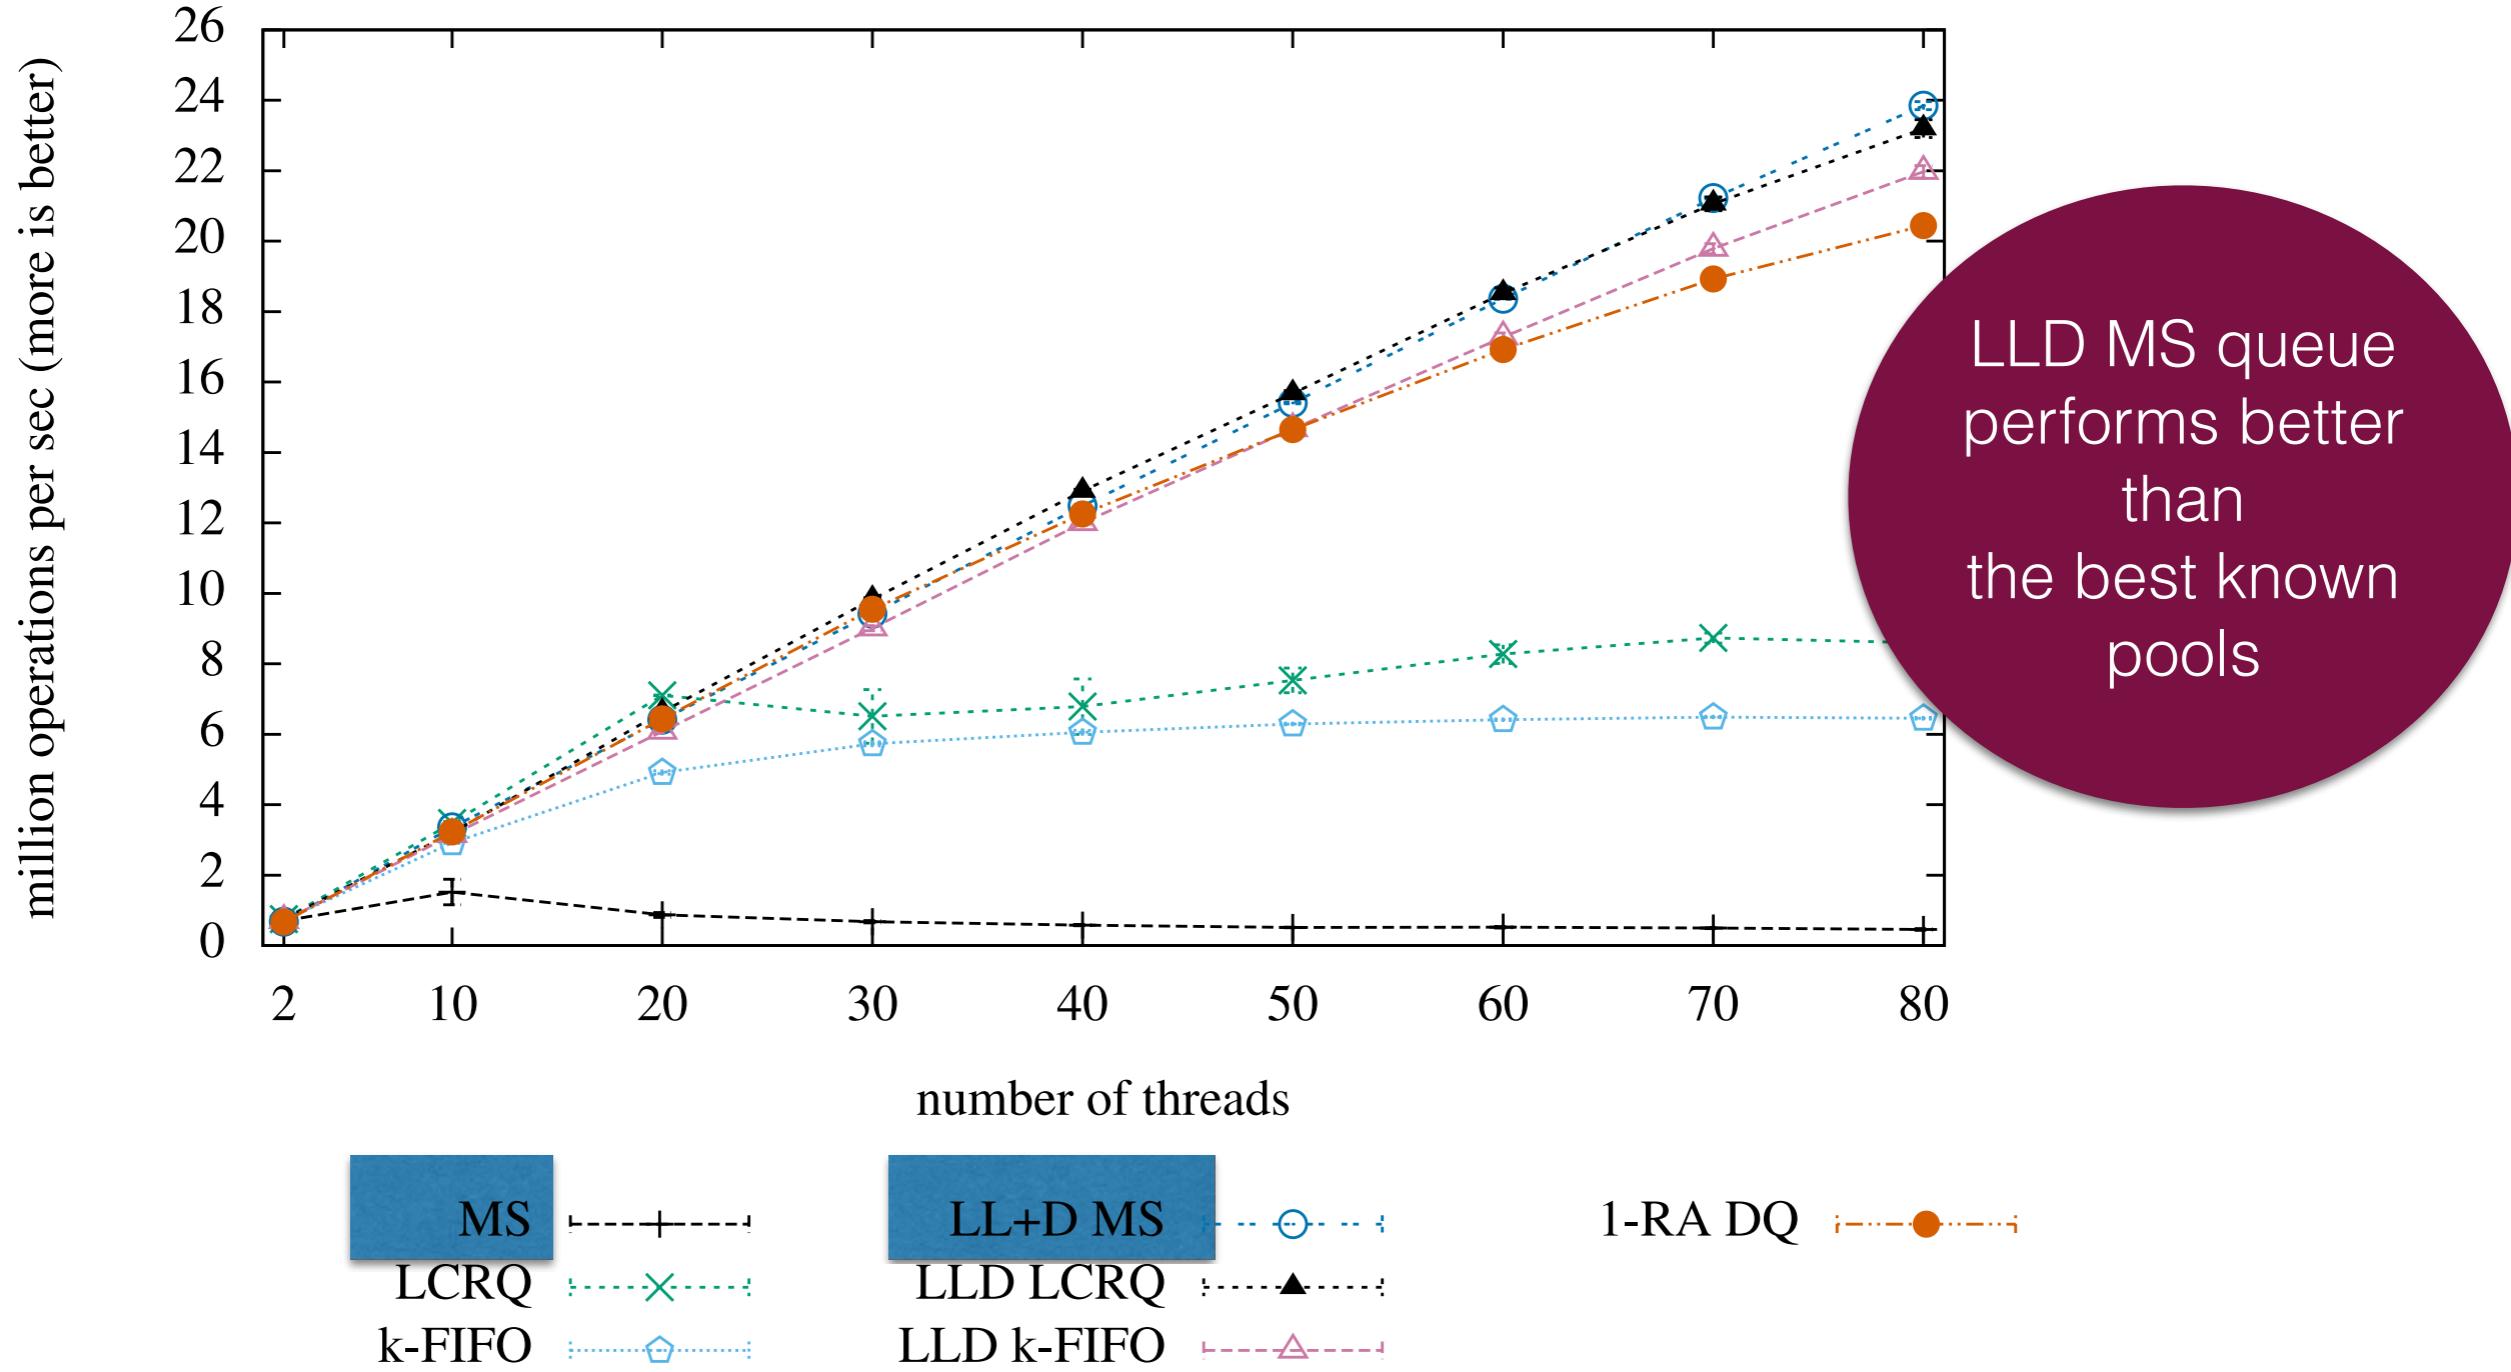
(a) Queues, LL queues, and “queue-like” pools

Performance



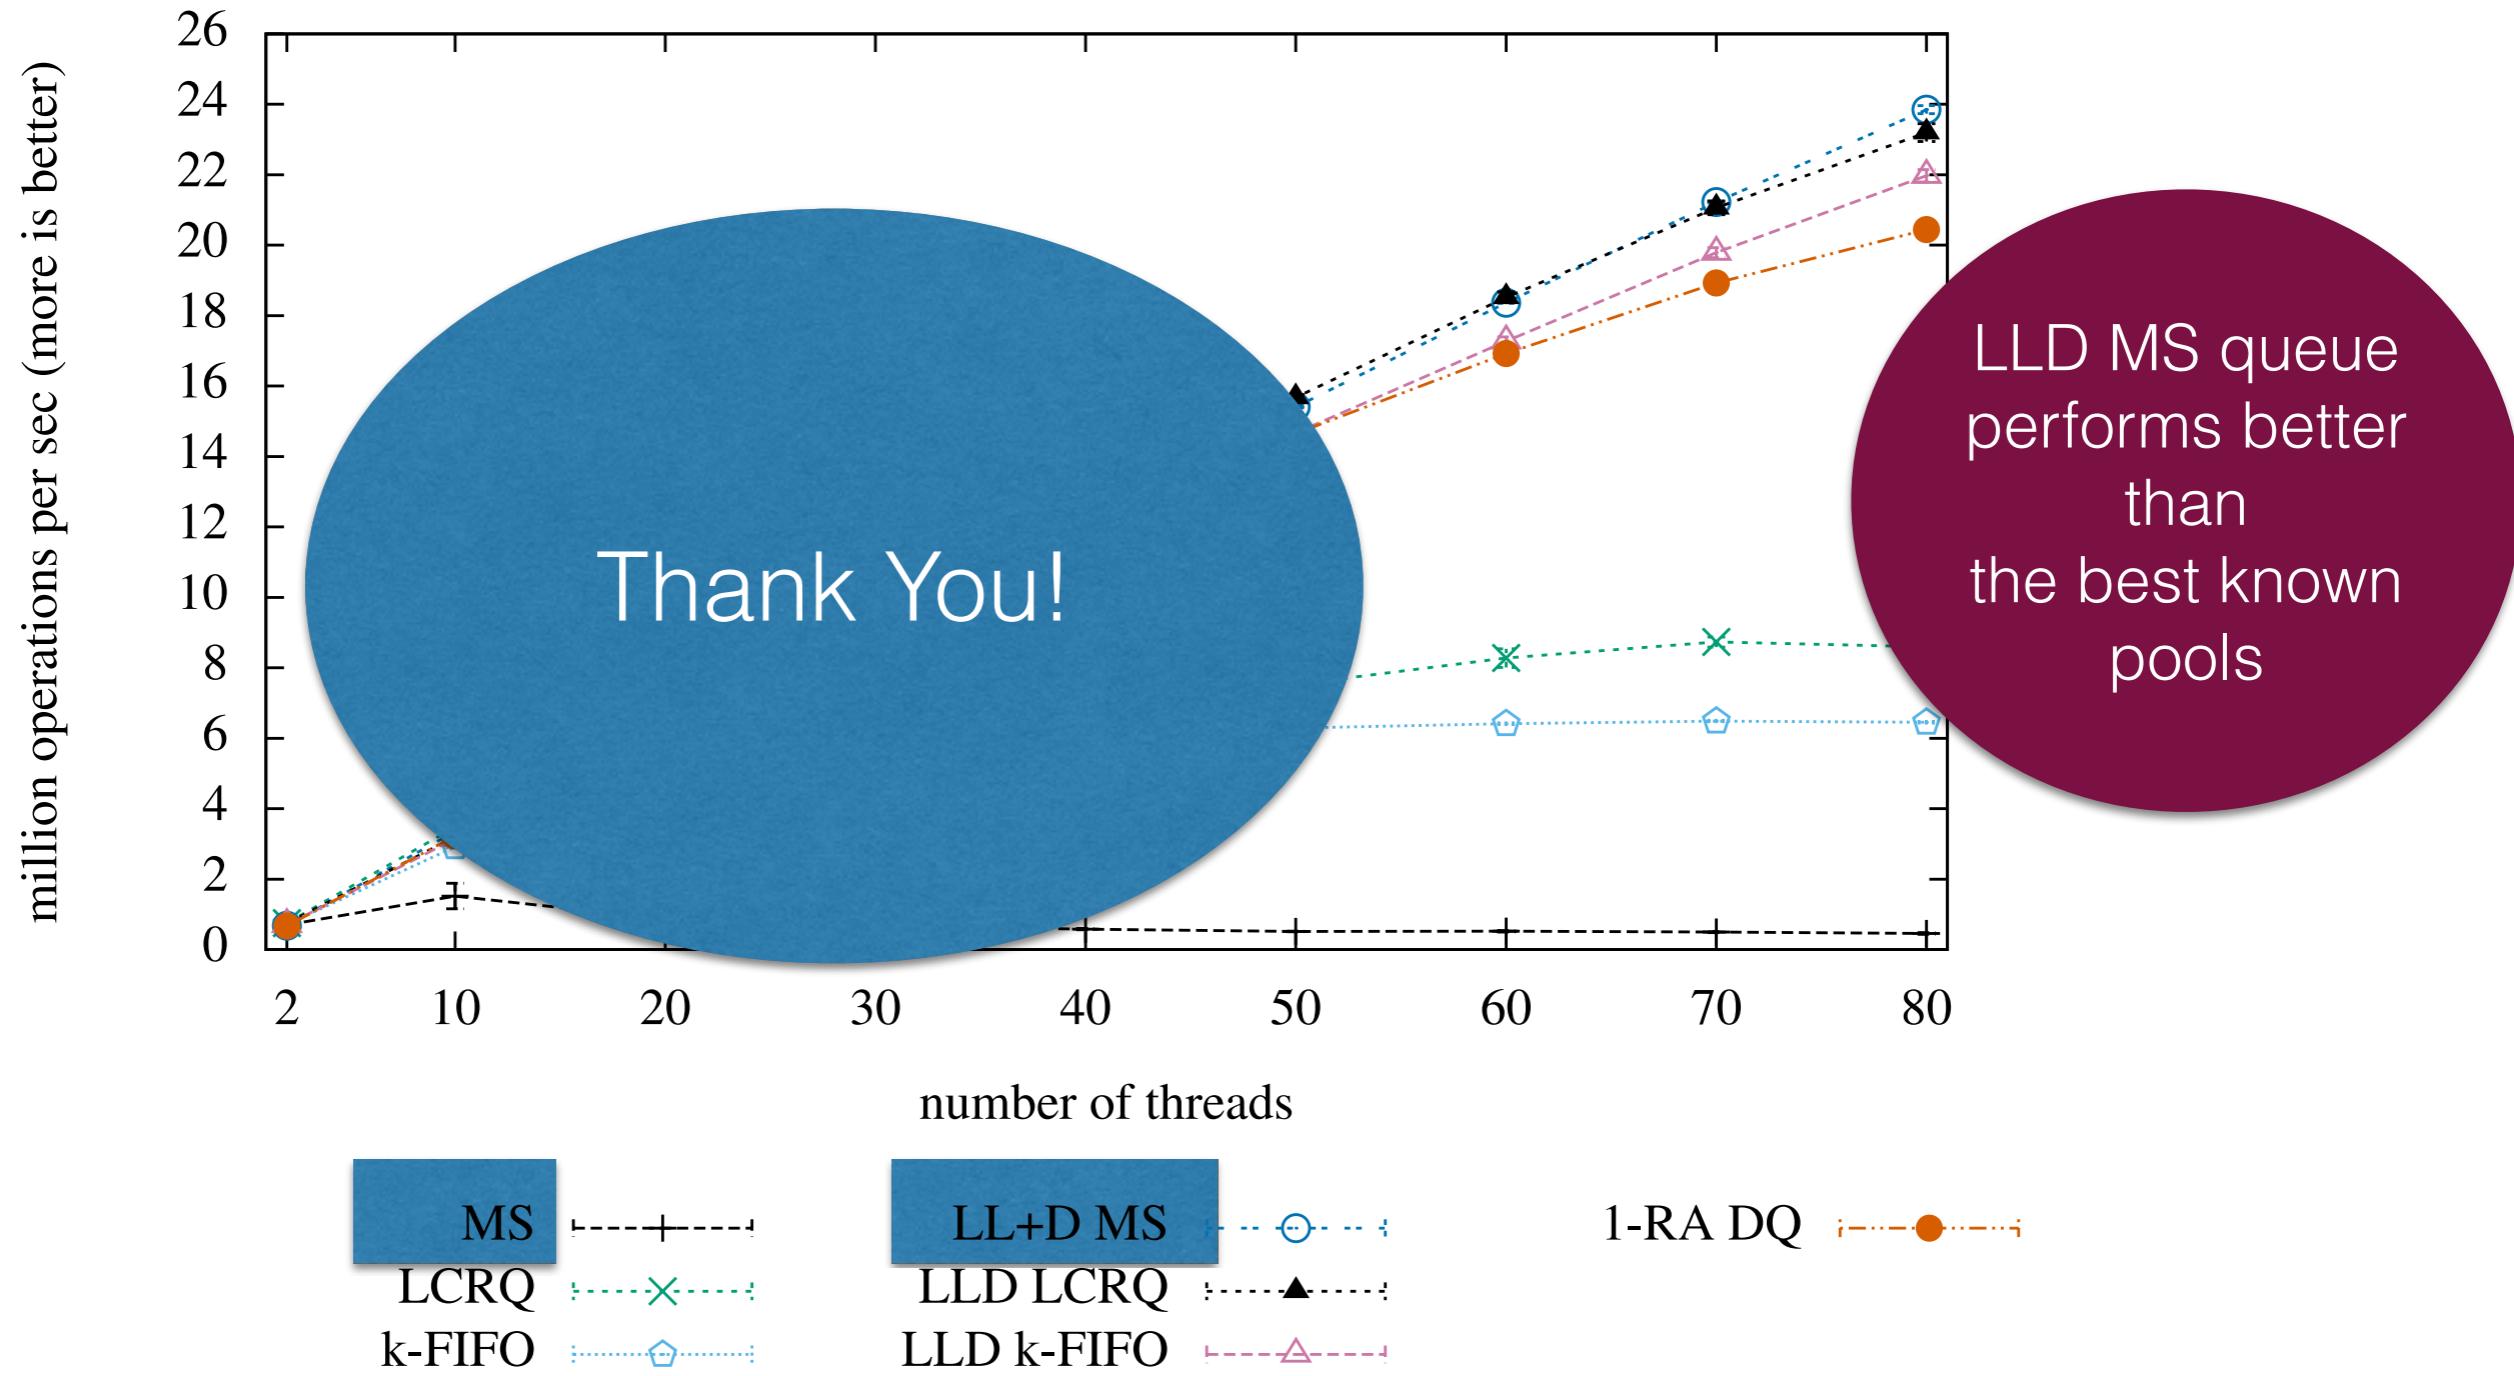
(a) Queues, LL queues, and “queue-like” pools

Performance



(a) Queues, LL queues, and “queue-like” pools

Performance



(a) Queues, LL queues, and “queue-like” pools